

Natural Language Processing
Episode 7 '2021


Model compression & acceleration

Yandex
Research




Chapter 1: why should you care?

Case study: text classification



Case study: text classification




Three kinds of efficiency:



Model Size
(mega)bytes

Case study: text classification



Three kinds of efficiency:




Model Size
(mega)bytes



Throughput
samples/second

Case study: text classification



Three kinds of efficiency:



Model Size
(mega)bytes




Throughput
samples/second

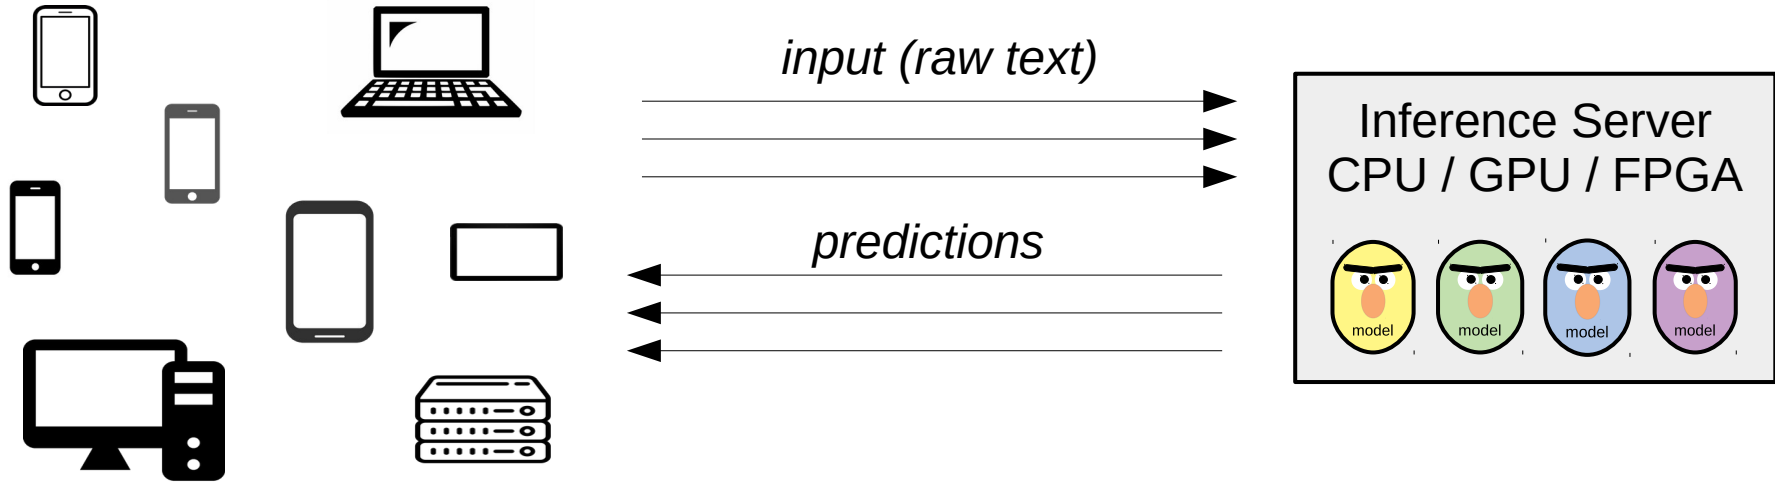


Latency
ms@percentile


Scenario 1: inference server



Scenario 1: inference server

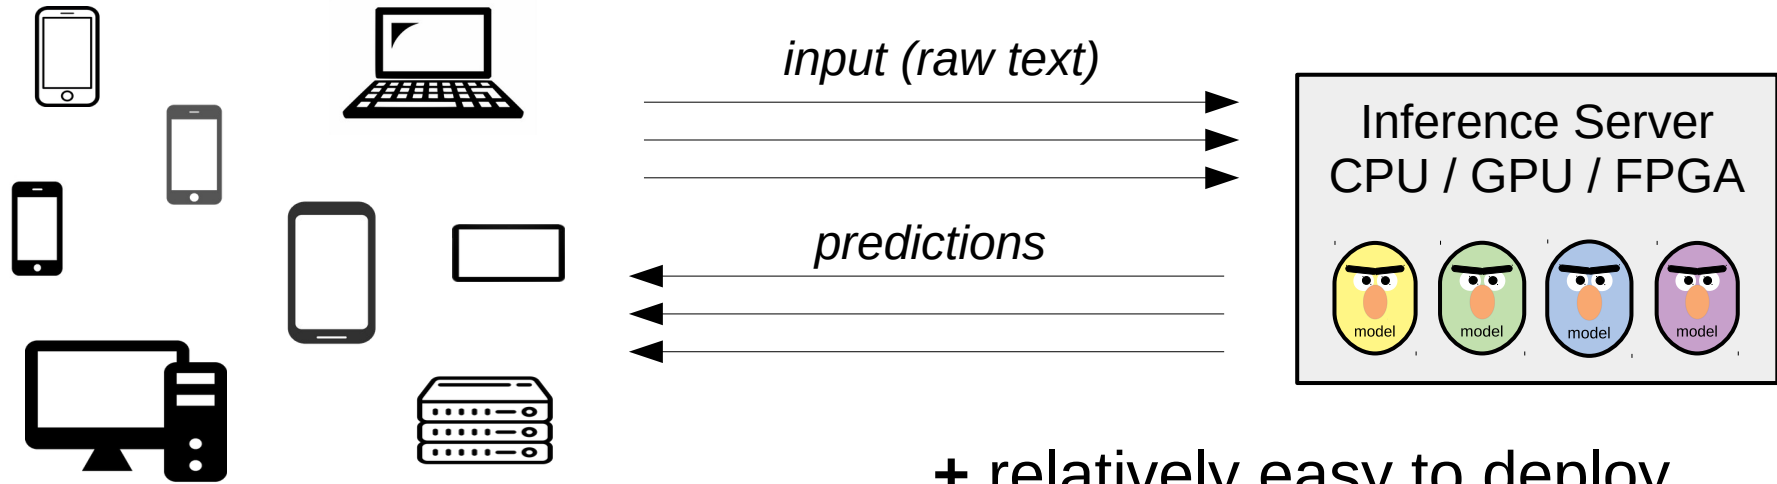


Scenario 1: inference server



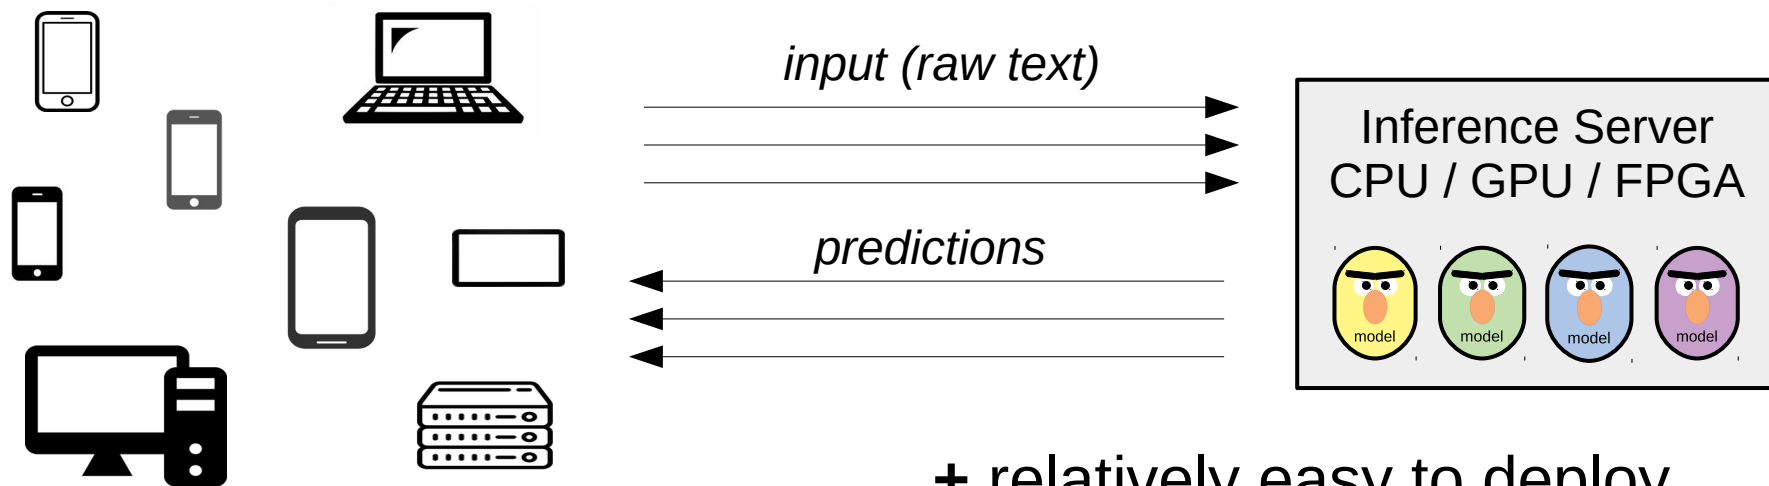
- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute

Scenario 1: inference server



- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute
- you pay for each inference
- clients can't work offline
- network latency

Scenario 1: inference server

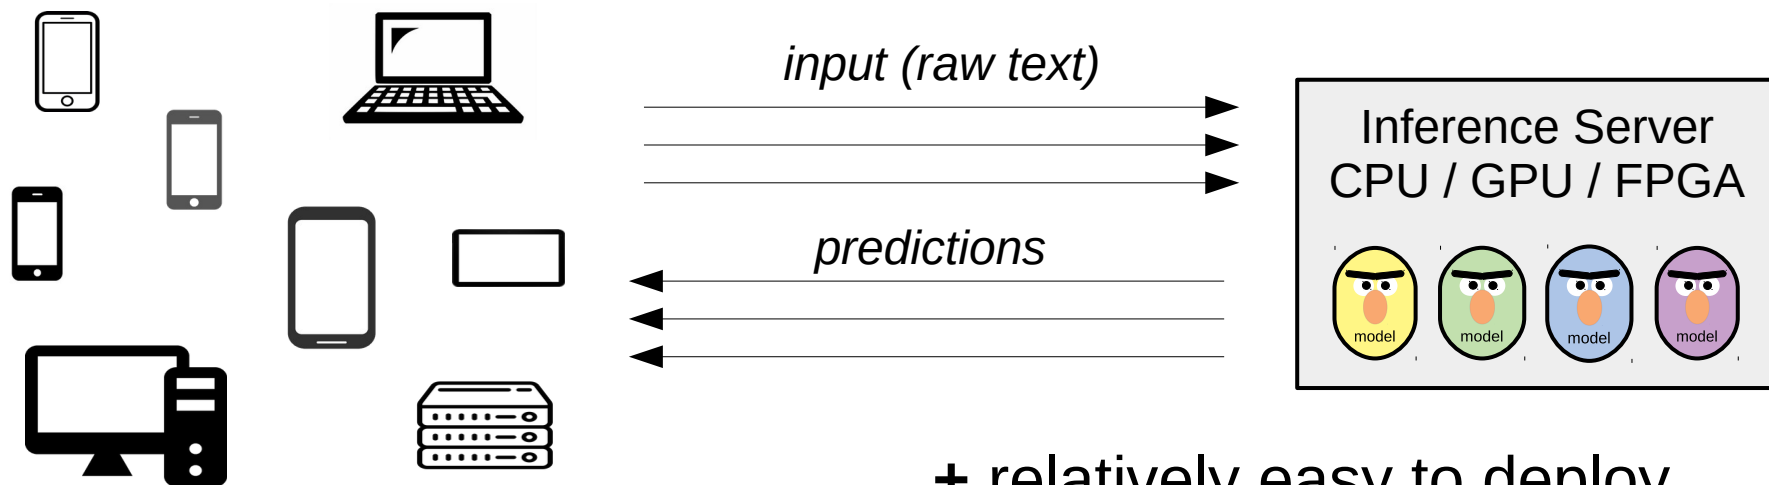


Which is the most important?



- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute
- you pay for each inference
- clients can't work offline
- network latency

Scenario 1: inference server



Priorities:



Note: smaller model = you can fit more models in the same memory

- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute
- you pay for each inference
- clients can't work offline
- network latency

Scenario 1: inference server

- Group inputs into batches (e.g. by length)
improves throughput at the cost of latency
- Multiple servers with load balancing
improves throughput at the cost of your budget :)

Scenario 1: inference server

- Group inputs into batches (e.g. by length)
improves throughput at the cost of latency
- Multiple servers with load balancing
improves throughput at the cost of your budget :)

Popular frameworks:

priorities



TensorFlow Serving

efficiency \ll developer time



TensorRT Inference Server (Triton)

efficiency \approx developer time



Custom model-dependent code

efficiency \gg developer time

Scenario 2: local inference

Preload model onto a dedicated device, infer locally using that device

Typical use cases:

- Parallel speech recognition
- “Smart” cameras
- Autonomous drones
- Self-driving cars


Priorities:



Scenario 3: web/smartphone app


- Load model weights on the fly and infer locally
Model size is critical for both you and the user

Scenario 3: web/smartphone app

- Load model weights on the fly and infer locally
Model size is critical for both you and the user
- Autonomous machine translation (tinyurl.com/yandex-translate-app)
- Pix2pix demo in a browser (<https://affinelayer.com/pixsrv>)
- Priorities: 

Scenario 3: web/smartphone app

- Load model weights on the fly and infer locally
Model size is critical for both you and the user
- Autonomous machine translation (tinyurl.com/yandex-translate-app)
- Pix2pix demo in a browser (<https://affinelayer.com/pixsrv>)

- Priorities:     

- Popular frameworks:



TensorFlow.js



CoreML



NNAPI

Platform

All modern browsers

iOS devices

Android devices


Chapter 2: how do I improve my model?

Compression by Distillation



Distillation...
Heard that word before?

Compression by Distillation

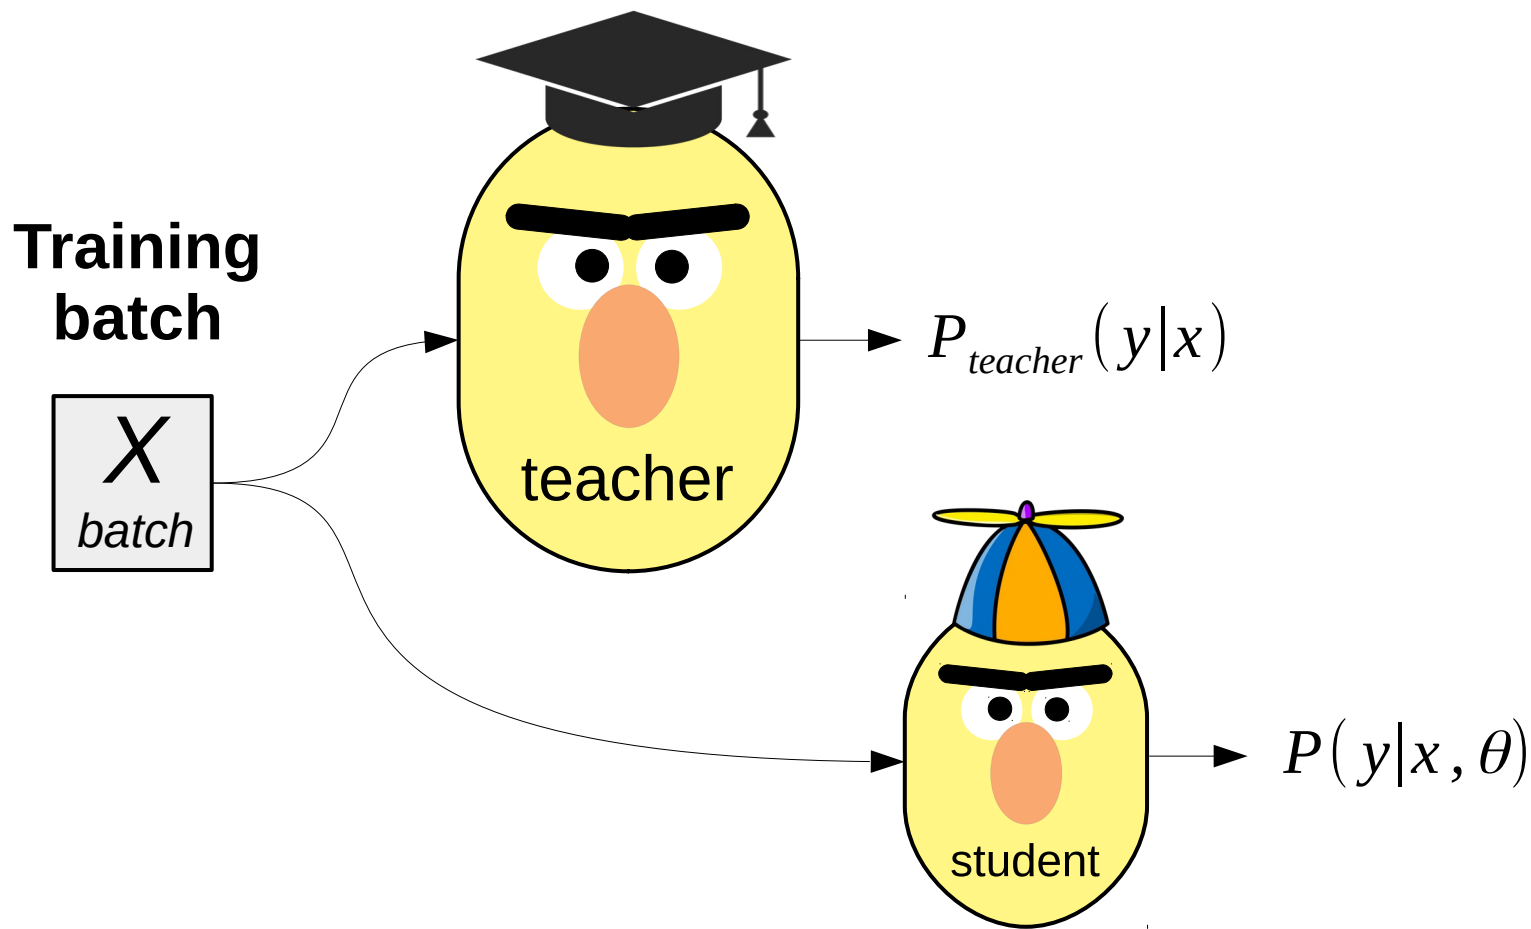


First, get the best performing model regardless of size

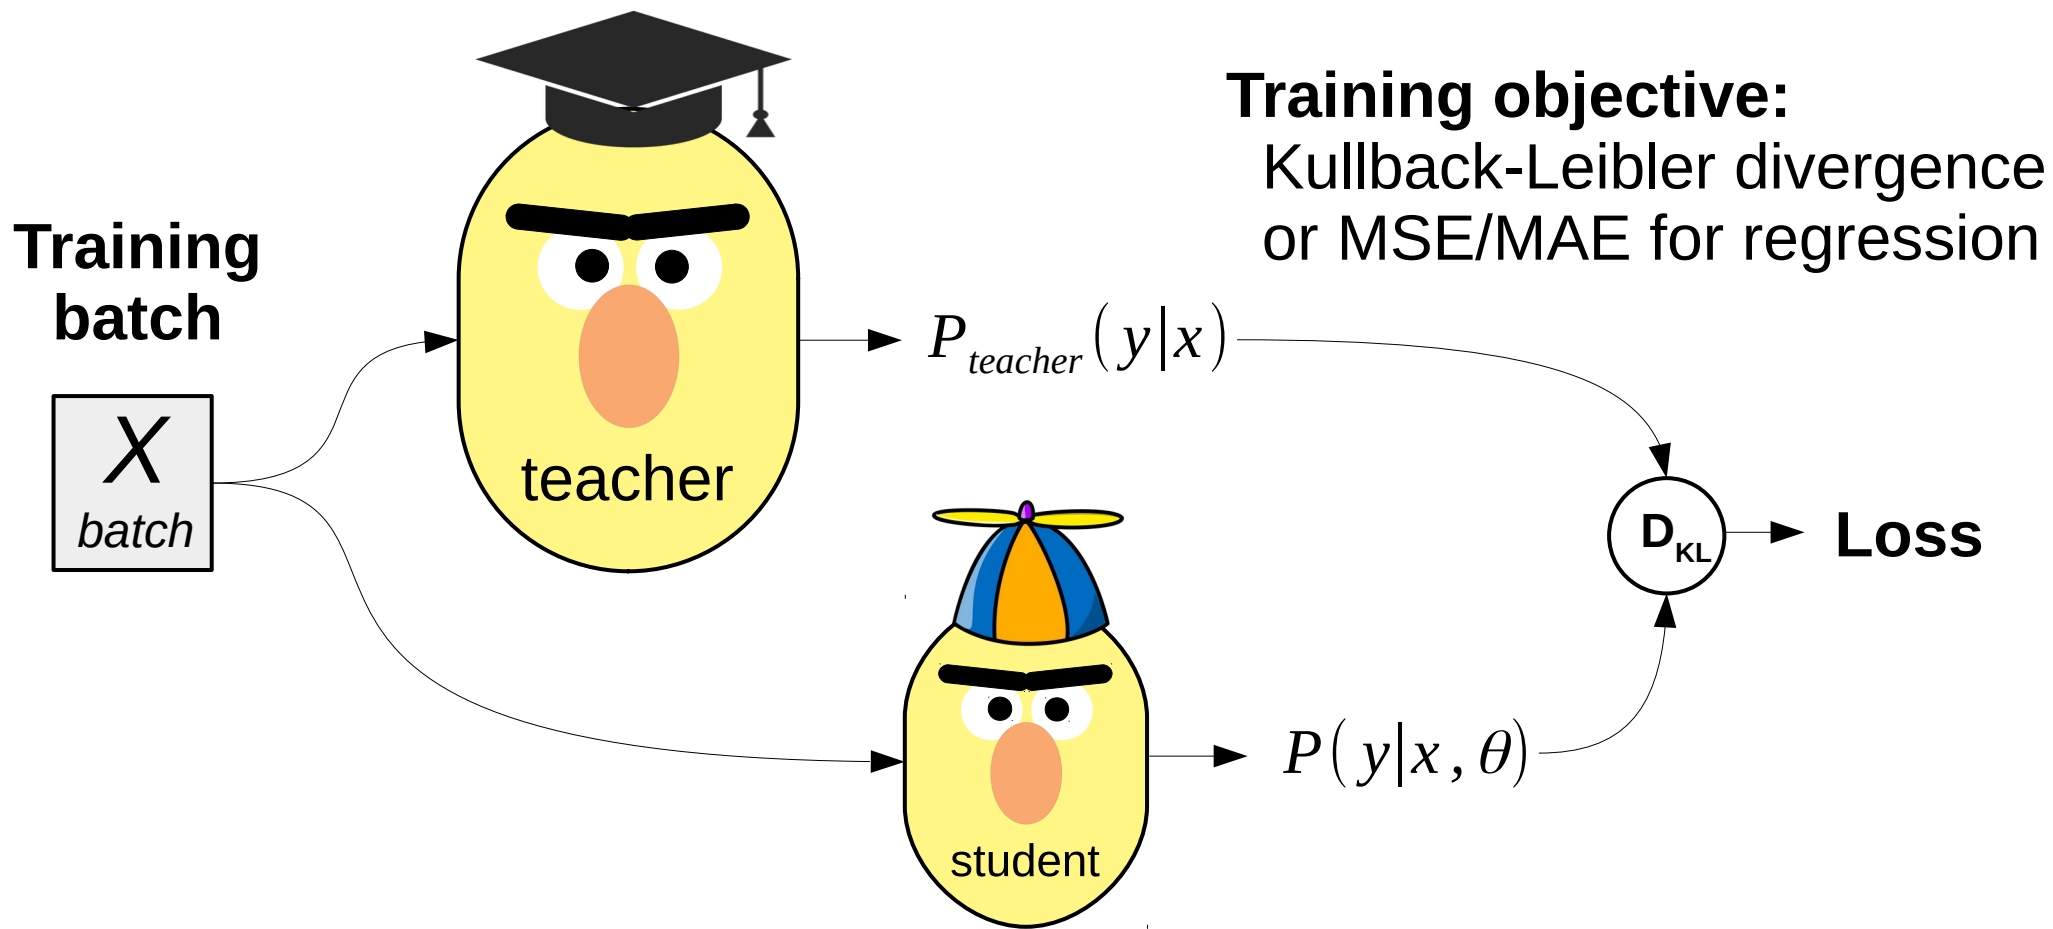


Then, train a more compact model to approximate it!

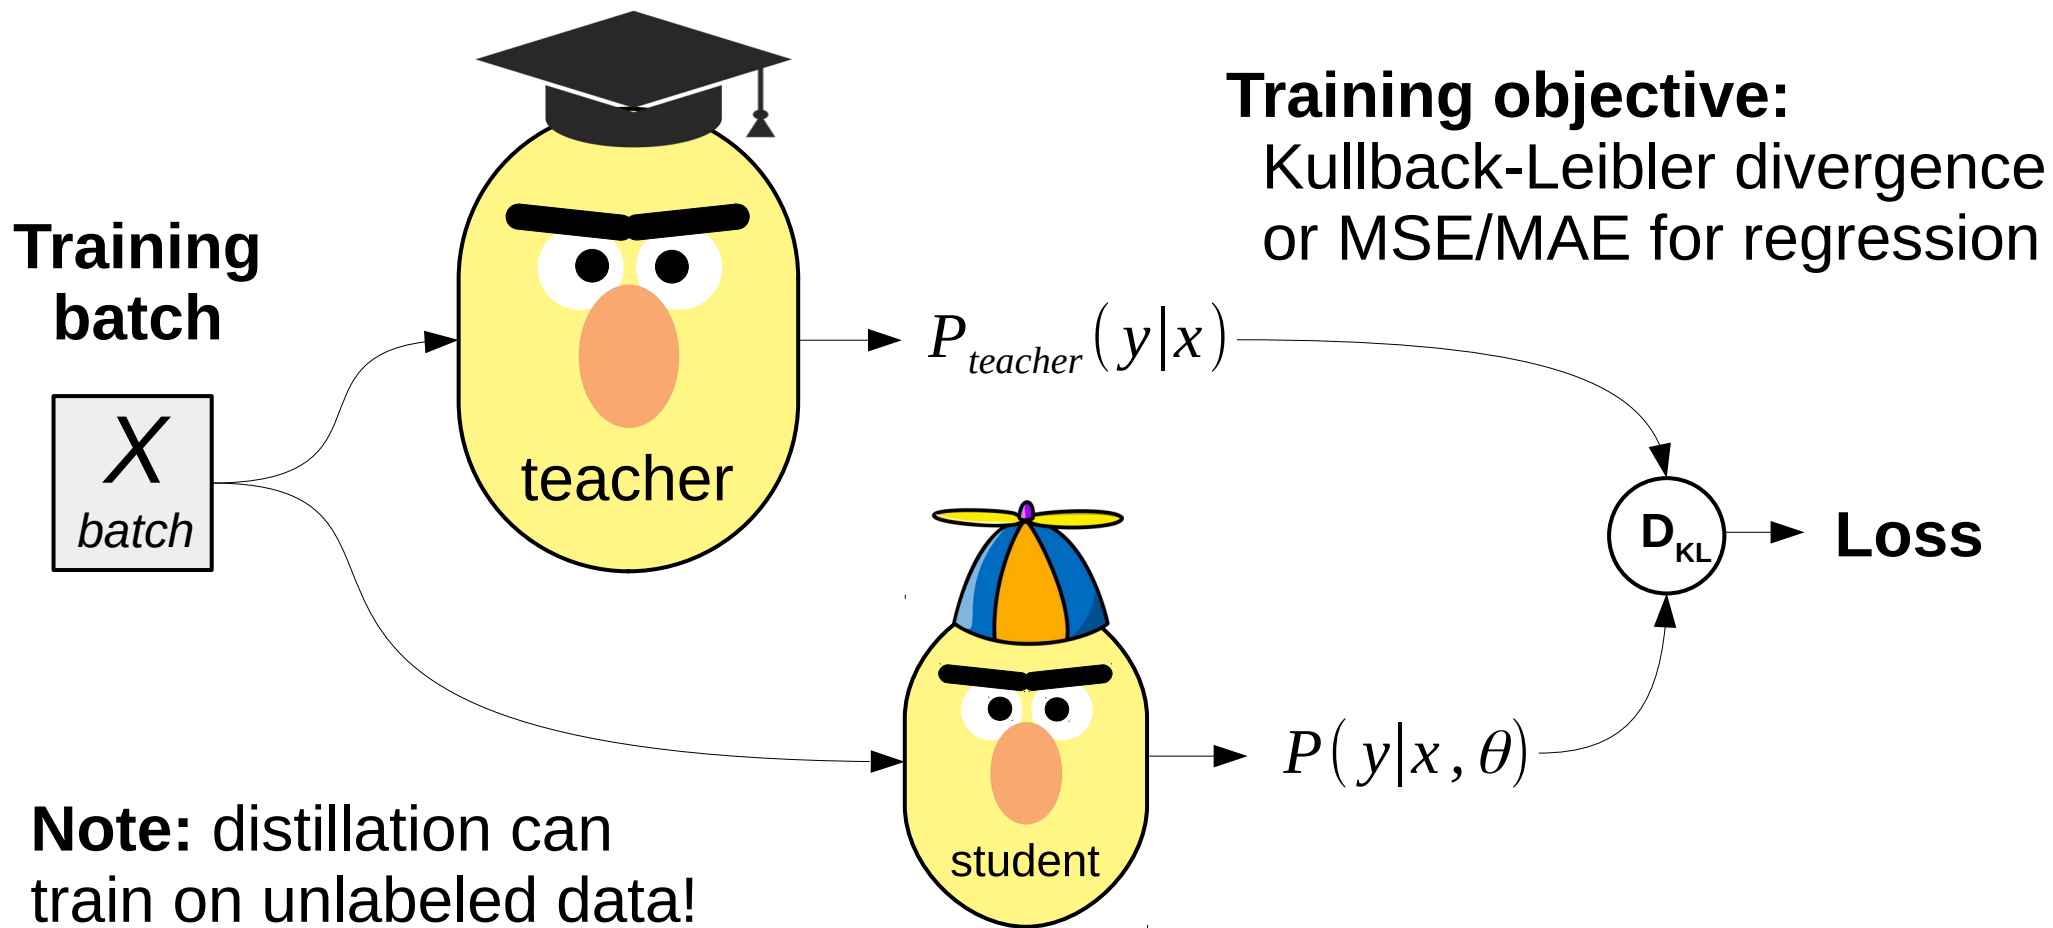
Compression by Distillation



Compression by Distillation



Compression by Distillation



Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

e.g. DistilBERT: <https://arxiv.org/pdf/1910.01108.pdf>

Same as BERT-base, but
with *half as many layers*
(and ≈ 1.5 times faster)

Model	# param. (Millions)	Inf. time (seconds)
ELMo	180	895
BERT-base	110	668
DistilBERT	66	410

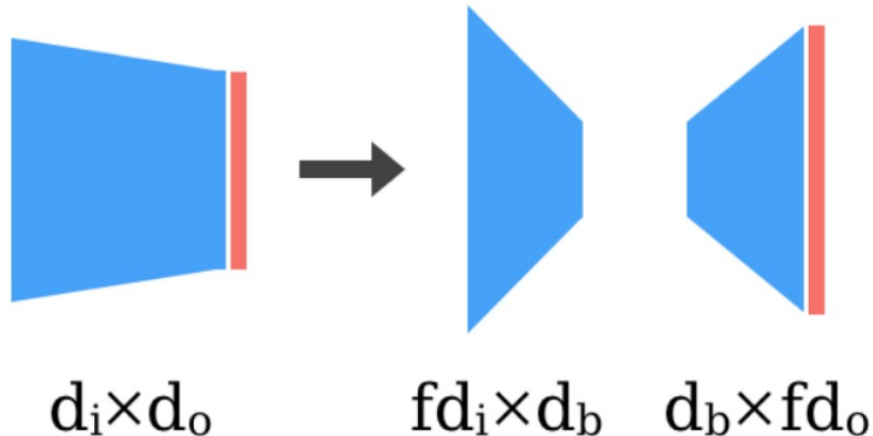
Model	Score	CoLA	MNLI	MRPC	QNLI	QQP	RTE	SST-2	STS-B	WNLI
ELMo	68.7	44.1	68.6	76.6	71.1	86.2	53.4	91.5	70.4	56.3
BERT-base	79.5	56.3	86.7	88.6	91.8	89.6	69.3	92.7	89.0	53.5
DistilBERT	77.0	51.3	82.2	87.5	89.2	88.5	59.9	91.3	86.9	56.3

Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

Factorized: product of smaller matrices or tensors




Images: https://openreview.net/pdf?id=_zx8Oka09eF

Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

Factorized: product of smaller matrices or tensors



Images: https://openreview.net/pdf?id=_zx8Oka09eF

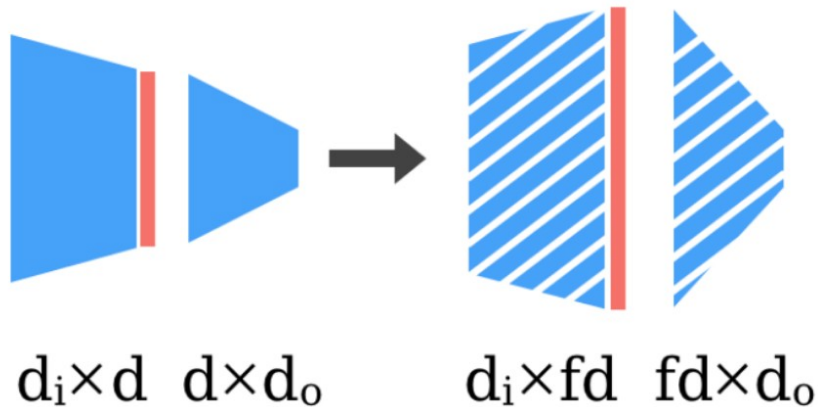
Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

Factorized: product of smaller matrices or tensors

Sparse: only a small (random) subset of weights are nonzero



Images: https://openreview.net/pdf?id=_zx8Oka09eF


Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

Factorized: product of smaller matrices or tensors

Sparse: only a small (random) subset of weights are nonzero



Q: how to store sparse weights?

Images: https://openreview.net/pdf?id=_zx8Oka09eF

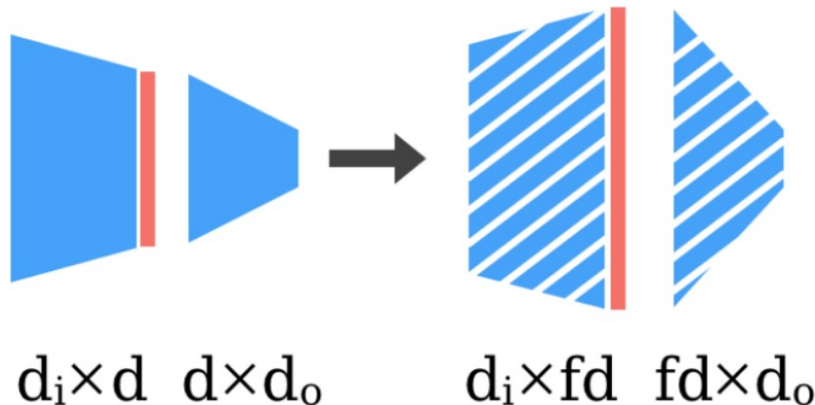
Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

Factorized: product of smaller matrices or tensors

Sparse: only a small (random) subset of weights are nonzero



Storage: only store random seed and nonzero weights.

Compute: sparse matrix multiply

Images: https://openreview.net/pdf?id=_zx8Oka09eF

Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

Factorized: product of smaller matrices or tensors

Sparse: only a small fraction of weights are nonzero

Read more: https://openreview.net/pdf?id=_zx8Oka09eF

Also: factorized embeddings <https://arxiv.org/abs/1901.10787>

Also also: small-world sparse weights graphs for RNNs

<https://tinyurl.com/openai-blocksparse>

Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

Factorized: product of smaller matrices or tensors

Sparse: only a small fraction of weights are nonzero

Read more: https://openreview.net/pdf?id=_zx8Oka09eF

Also: factorized embeddings <https://arxiv.org/abs/1901.10787>

Also also: <https://tinyurl.com/openai-blocksparse>




- More distillation tricks:

Ensemble distillation <https://arxiv.org/abs/1702.01802>

Dropout distillation <http://proceedings.mlr.press/v48/bulo16.pdf>

Co-distillation <https://arxiv.org/abs/1804.03235>

Compression by quantization



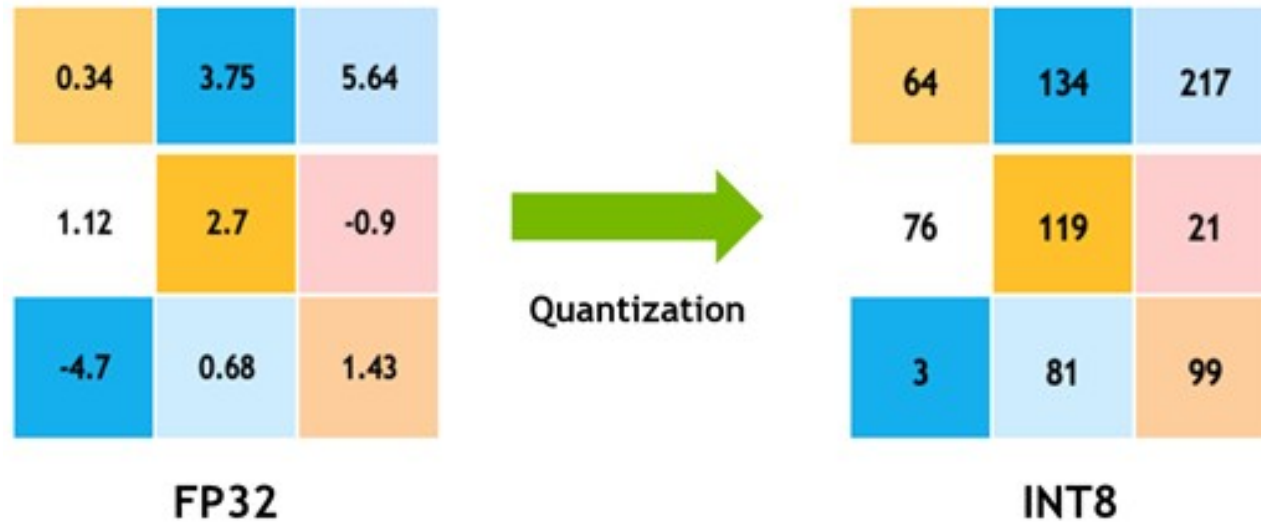
FP16



INT8




Linear quantization




- 1) Scale inputs and parameters into uint8 range
- 2) Multiply in uint8, accumulate to int32
- 3) Un-scale multiplication results in float32

Non-linear quantization




Consider weights as a distribution

Non-linear quantization




Consider weights as a distribution

Non-linear quantization




Consider weights as a distribution

Non-linear quantization



Compute a grid of percentiles

Non-linear quantization



percentiles (32-bit)


Index (4- or 8-bit) of
nearest percentile
for each weight

Store each weight as its nearest percentile


High-dimensional case

Quantize entire vectors as K-means

Quantization Example



```
quantizer = KMeans(n_clusters=7).fit(X)
```



Images: [Jeremy Jordan](#)

OPQ, AQ, LSQ

Product Quantization

Split vectors into chunks, quantize each chunk separately

Orthogonal Product Quantization

First run orthogonal transform, then product quantization

<http://kaiminghe.com/publications/cvpr13opq.pdf>

More:

Additive Quantization

<https://tinyurl.com/babenko-aq-pdf>

Local Search Quantization

<https://tinyurl.com/martinez-lsq-pdf>

Low-precision training!

Training with 8-bit params:

<https://arxiv.org/abs/1812.08011>

<https://arxiv.org/abs/1805.11046>

Compressing gradients to 1 bit and beyond

<https://arxiv.org/abs/2102.02888>

<https://arxiv.org/abs/1905.13727>

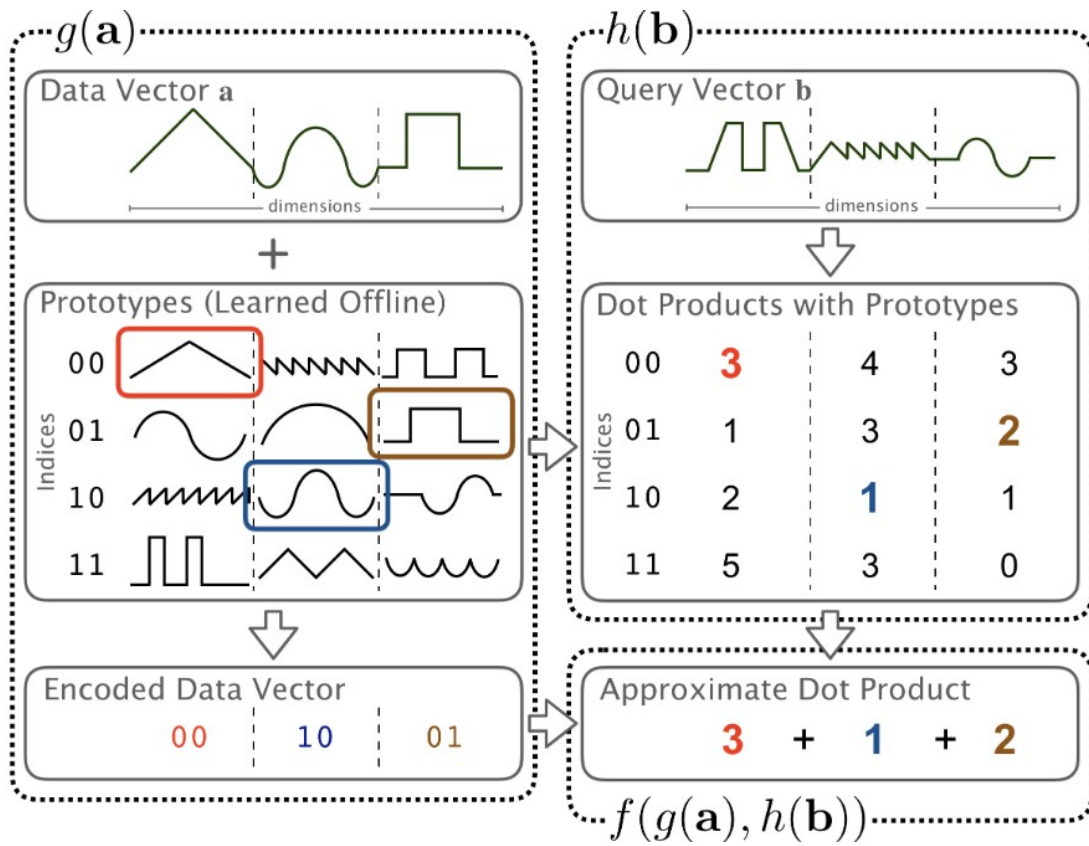
Training with 8-bit optimizers:

<https://www.youtube.com/watch?v=IxrlHAJtqKE>

<https://arxiv.org/abs/2110.02861>


Quantization + pre-computation

Source: <https://arxiv.org/abs/2106.10860>




- product-quantize weights
- “quantize” inputs via LSH
- pre-compute dot products all inputs x all codes
- “multiply” by **looking up** pre-computed products

Acceleration by cascading

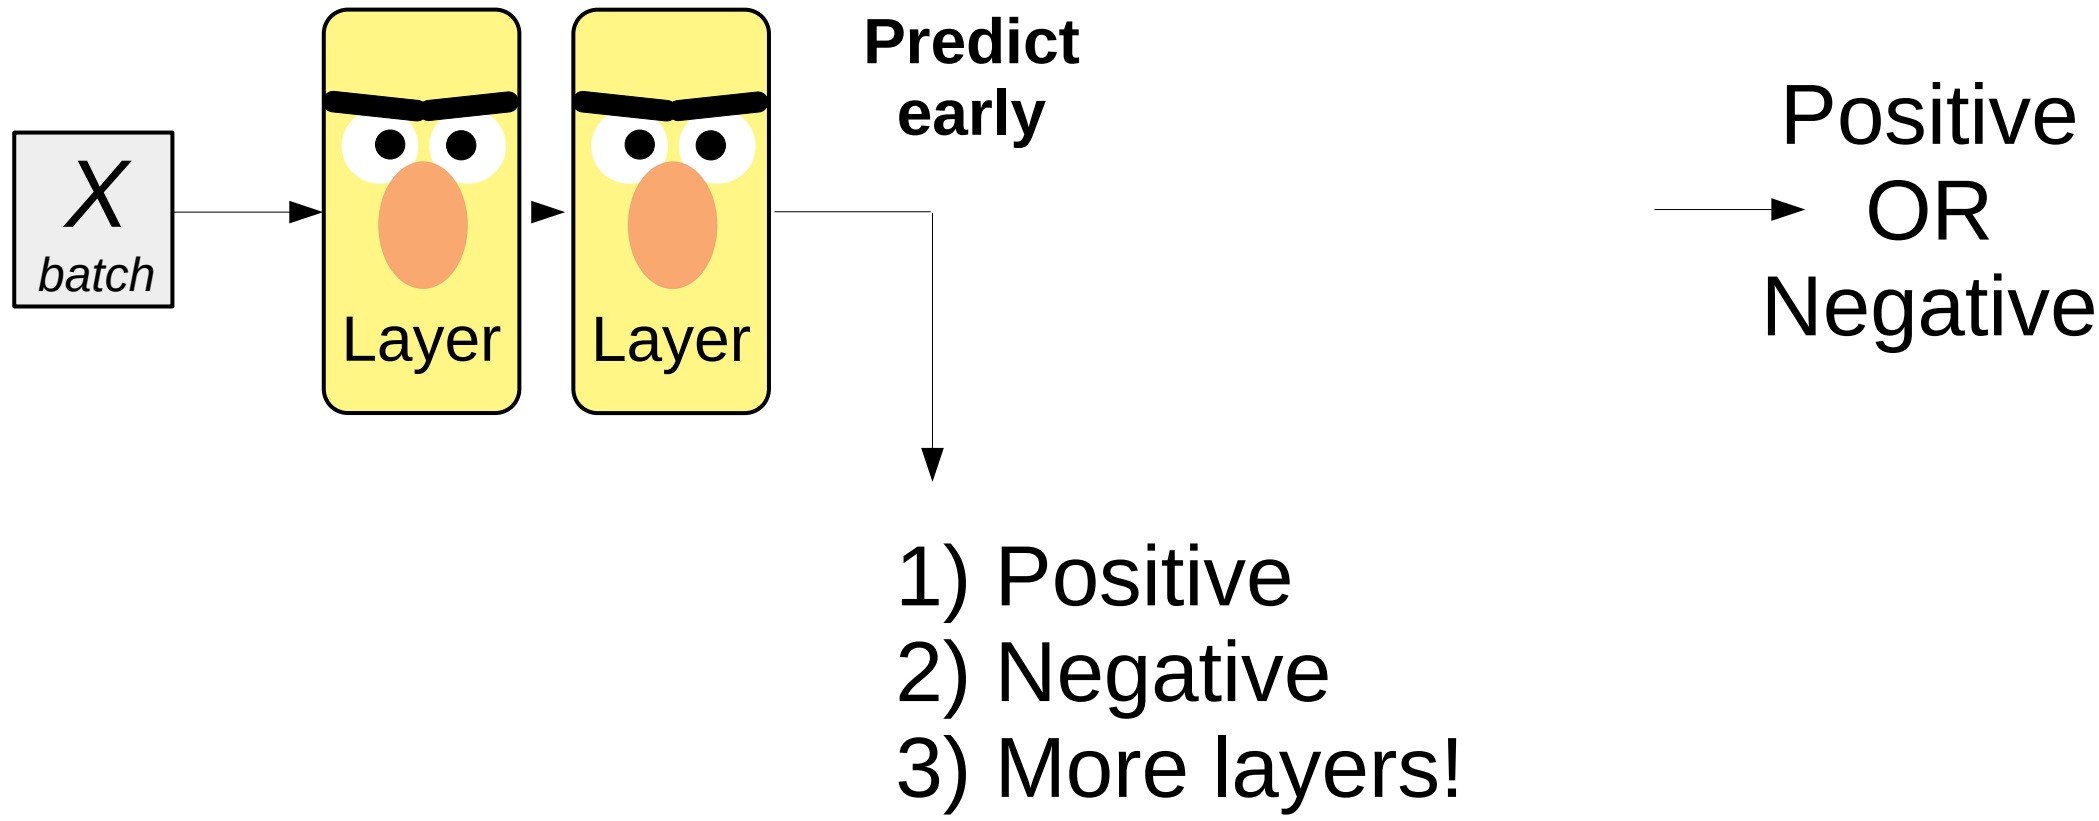


Acceleration by cascading




Do we really need every layer
all the time?

Acceleration by cascading




Acceleration by cascading

Only if “more layers”



- 1) Positive
- 2) Negative
- 3) More layers!

Adaptive Computation Time



Original ACTI (for RNN)
<https://arxiv.org/abs/1603.08983>

Spatial ACT (conv)
<https://tinyurl.com/sact-pdf>


ACT Transformers
<https://arxiv.org/abs/1807.03819>

Compression by sparsification

Do we really need all D by D weights?

Compression by pruning

Do we really need all D by D weights?



Magnitude pruning

Drop ~5% smallest weights
from each layer every 1000 steps
(and keep training)

Reminds you of something?

Magnitude pruning

Drop ~5% smallest weights
from each layer every 1000 steps
(and keep training)

Reminds you of something?
See ML course, Optimal Brain Damage

Pruning with L_0 regularization


Add a special regularizer that encourages dropping unnecessary weights

Whiteboard time!

Read more: <https://arxiv.org/abs/1712.01312>

Alternative: <https://arxiv.org/abs/1701.05369>

Which one works best?



Source <https://arxiv.org/abs/1902.09574>

Pruning with L_0 regularization

Add a special regularizer that encourages dropping unnecessary weights

Whiteboard time!

Pruning with L_0 regularization

Add a special regularizer that encourages dropping unnecessary weights


Can prune

- individual weights
- Individual neurons
- attention heads
- entire layers!


$$\lambda = 0.01$$

Pruning heads: https://lena-voita.github.io/posts/acl19_heads.html


Sparse training



Random




1-Dimensional Watts-Strogatz



Barabasi-Albert

<https://cdn.openai.com/blocksparse/blocksparspaper.pdf>

Sparse training



<https://cdn.openai.com/blocksparse/blocksparsenpaper.pdf>

What did we learn?