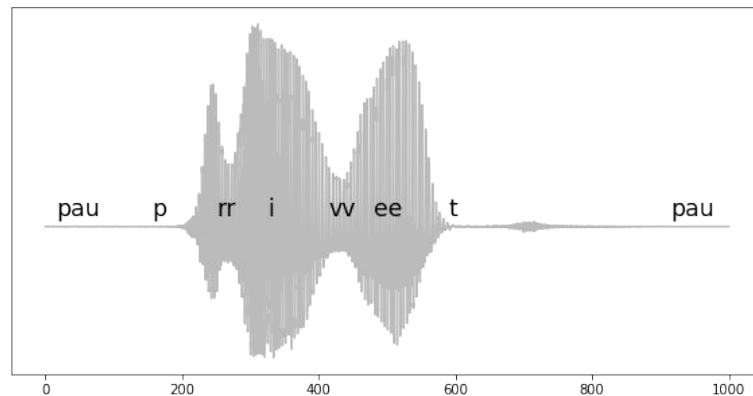



# Text-to-speech synthesis

Vladimir Kirichenko, Yandex

# Problem Definition

“Привет!”



# Problem Definition

Input:

Raw text (with digits, strange symbols, abbreviations etc.)

Output:

PCM audio (wav)



# Quality Assessment

- No “right” or “wrong” output.
- Subjective perception of the quality.
- Different kinds of errors

# Mean Opinion Score

1. Remove “hard” errors (words mispronounced, wrong intonation)
2. Ask assessors to rate each example on a scale of 1 to 5
3. Average the results

# Mean Opinion Score

**Table 1.** MOS (ACR) scores

Rating	Quality	Distortion
5	Excellent	Imperceptible
4	Good	Just perceptible, but not annoying
3	Fair	Perceptible and slightly annoying
2	Poor	Annoying, but not objectionable
1	Bad	Very annoying and objectionable

# Mean Opinion Score - CrowdMOS

- Scores determination and screening
- CI-s estimation
- Crowdsourcing MOS estimation with Amazon MTurk

# Mean Opinion Score

## Pros:

- Absolute scale
- Good for sound quality estimation

## Cons:

- Scores depends a lot on the crowdsourcing platform
- “Hard” errors are not considered
- Bad for intonation/pronunciation assessment



# MUSHRA (MUltiple Stimuli with Hidden Reference and Anchor)

1. For each estimated example provide “reference” with natural speech
2. Ask assessors how similar are reference and example sentence (on a scale of 1 to 5)
3. Average the results

# MUSHRA

## Pros:

- Absolute scale
- More stable than MOS
- Good for intonation and speaker similarity assessment

## Cons:

- Depends a lot on a test set
- Expensive to change test set
- Pronunciation issues aren't considered

# (pronunciation) Sentence Error Rate

1. Ask assessors to mark sentences with any “hard” errors
2. Average the results.

# (pronunciation) Sentence Error Rate

## Pros:

- Considers all “hard” errors
- Cheap collection of test set
- Good to estimate rate of pronunciation errors

## Cons:

- Intonation errors rate is very noisy
- Not sensitive to sound quality issues

# Side by Side Test

1. Get pairs of sentences (with the same phrase) from two different syntheses
2. Ask assessors to choose the most preferable audio for each pair
3. Average the votes for each synthesis

# Side by Side Test

## Pros:

- Sensitive to all kinds of issues
- Binomial test could help to estimate confidence of the SbS measurement
- Good to compare syntheses quality


## Cons:

- Relative scale
- Noisy in the case of almost-equal syntheses

# Datasets

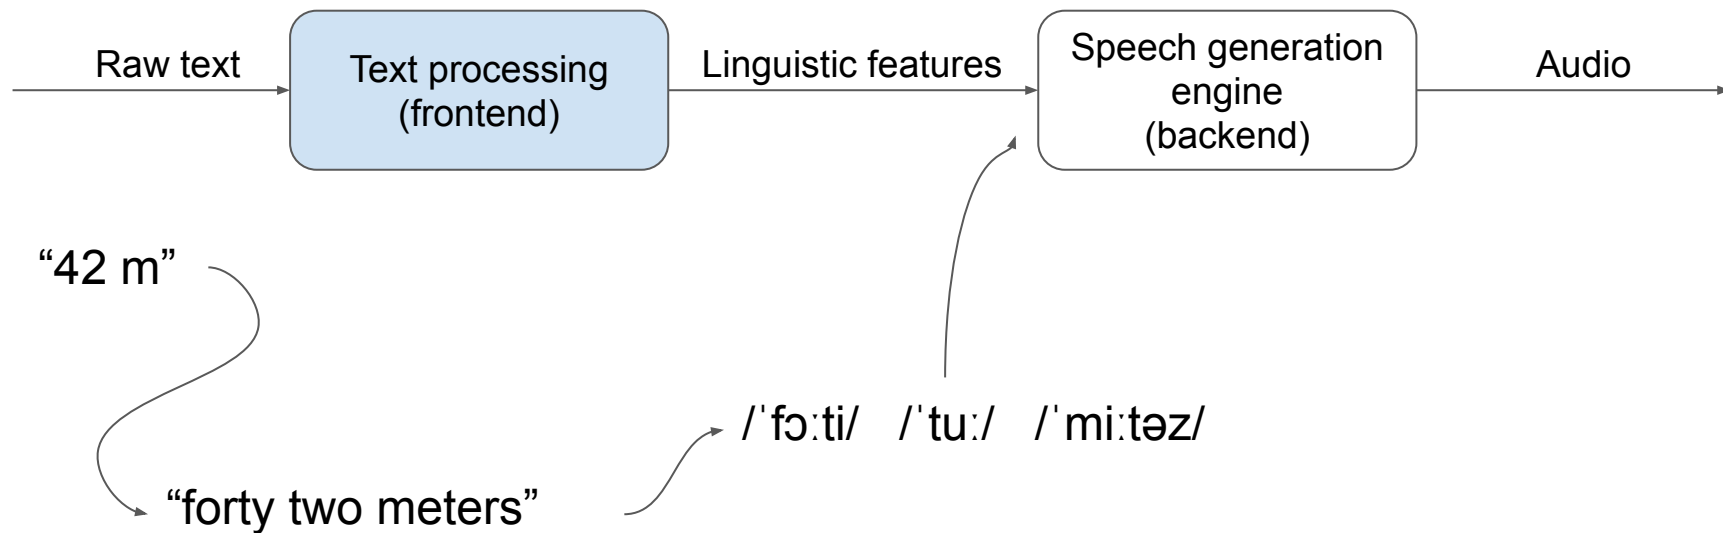
- LJ Speech
  - <https://keithito.com/LJ-Speech-Dataset/>
  - ~24h of 1 speaker (EN)
- VCTK
  - <https://homepages.inf.ed.ac.uk/jyamagis/page3/page58/page58.html>
  - ~44h of 109 speakers (EN)
- M-AILABS
  - <https://www.caito.de/2019/01/the-m-ailabs-speech-dataset/>
  - ~1000h, 8 languages, 16kHz
- Closed companies datasets

# TTS Pipeline

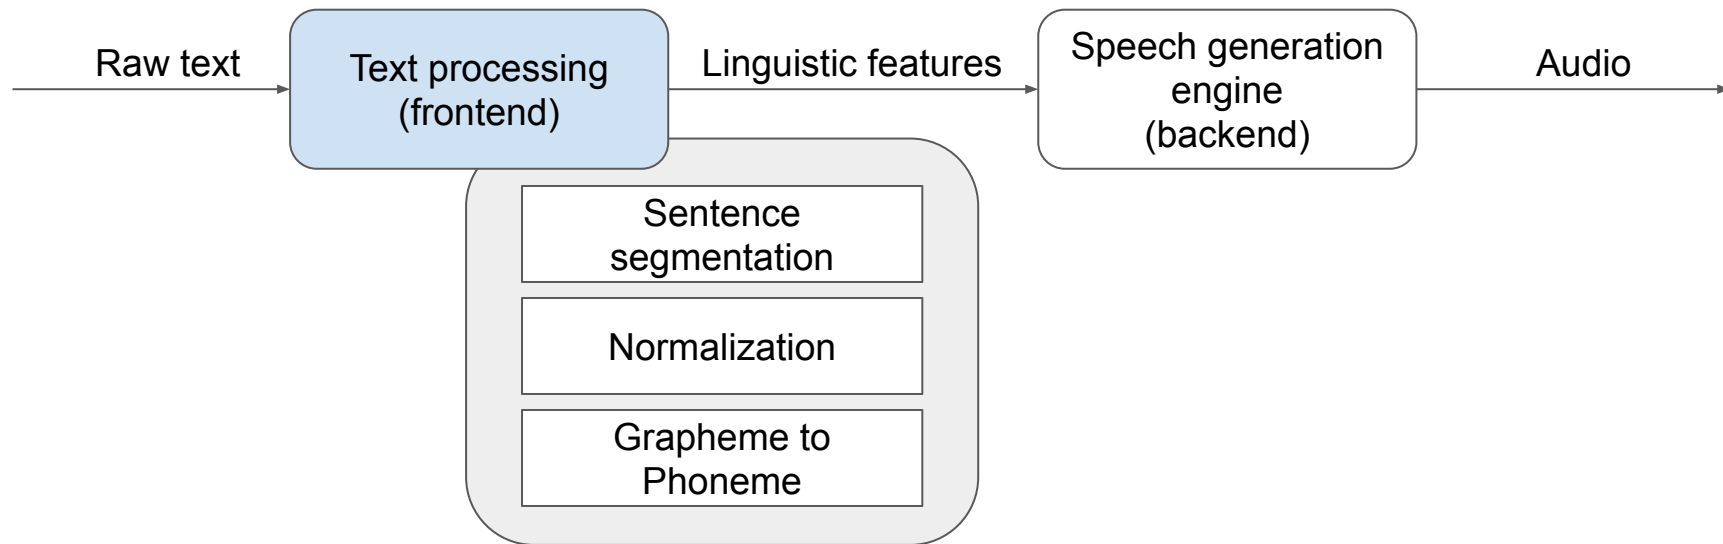




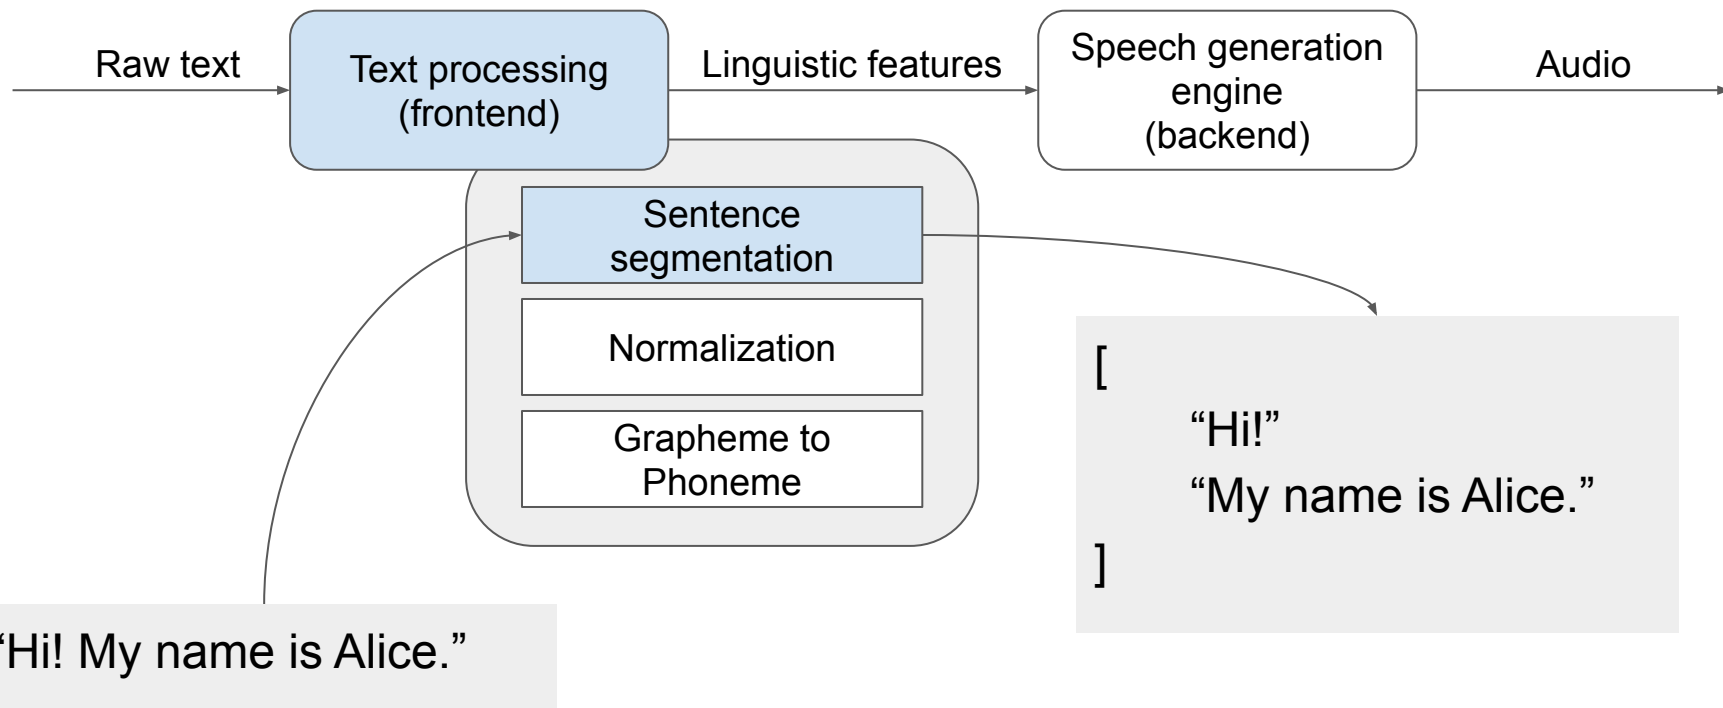
# Preprocessor



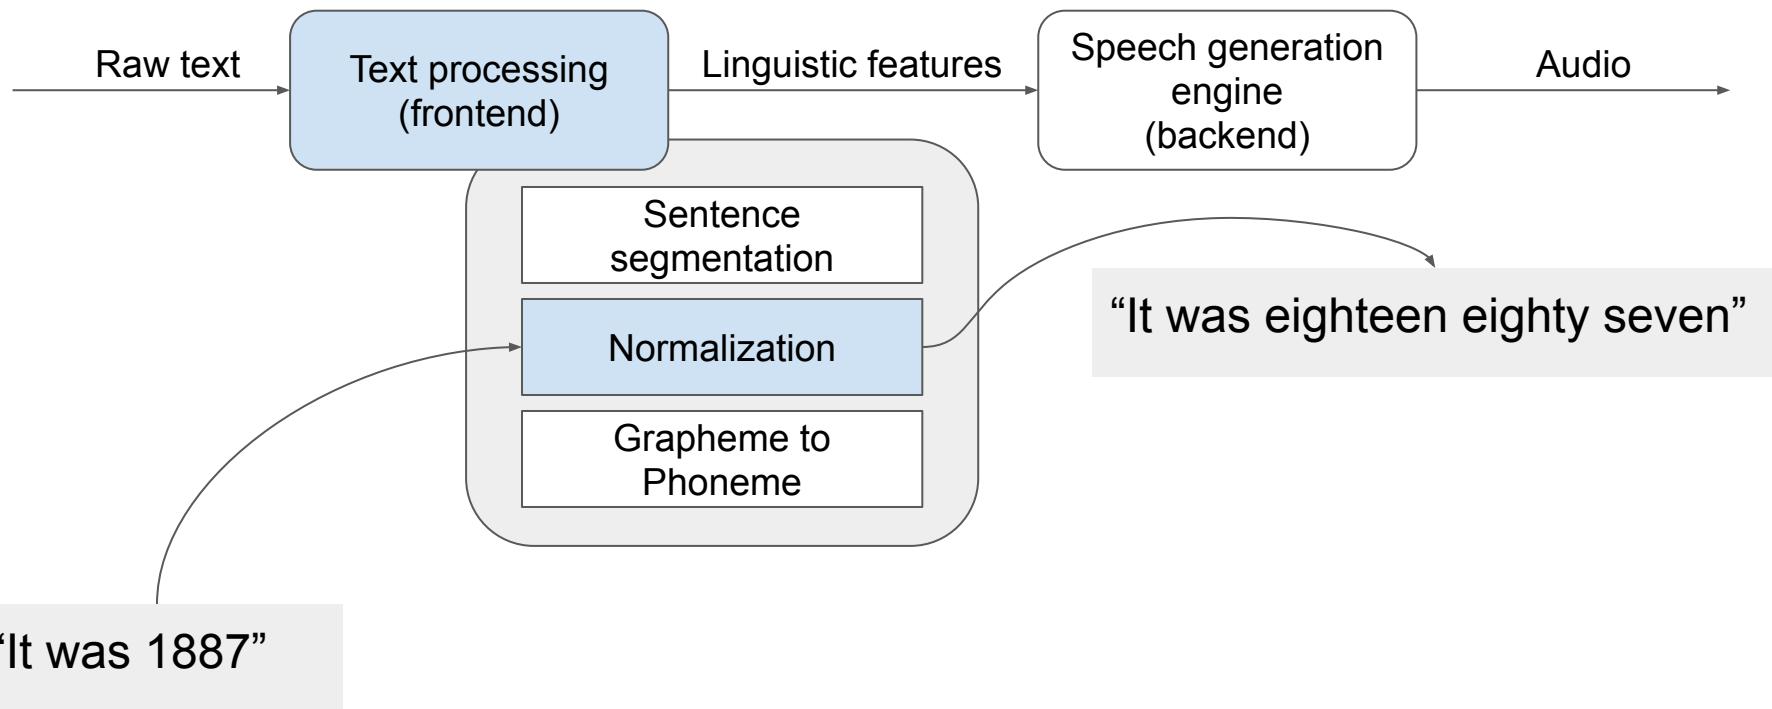
# Preprocessor



# Sentence tokenization




# Normalizer



# WFST Normalizer

Weighted Final State Transducer



# WFST Normalizer

Rules are weighted substitutions:

- `s/\b1\b/one/w=0.5` - replace “1” with “one”
- `s/\b1\s*st\b/first/w=1.0` - if it is not followed by “st”

# WFST Normalizer

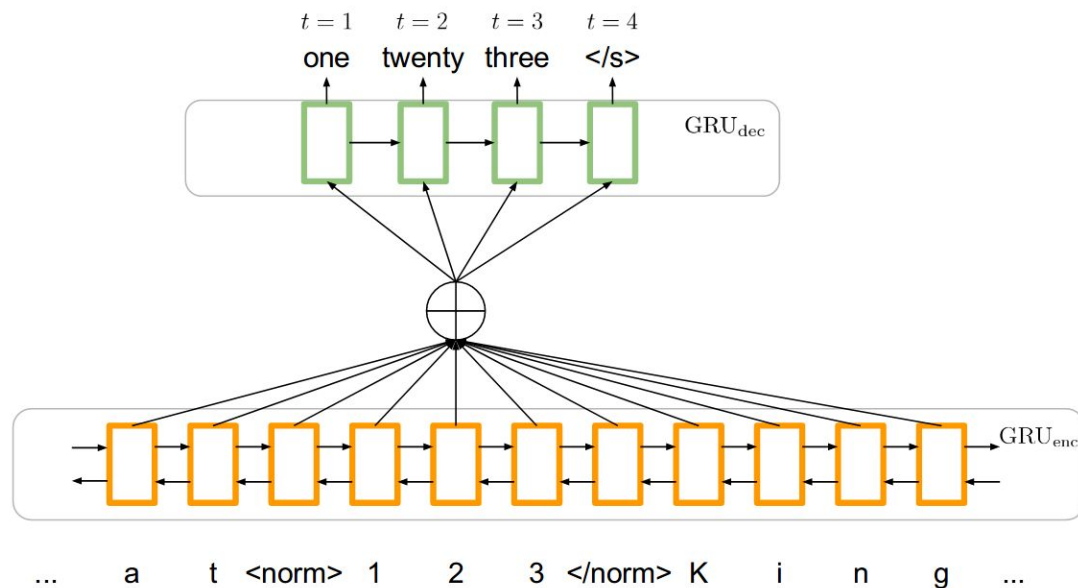
Context problem:

- “Дом 3 к. 3” - “дом три кубические копейки”
- “Нужно 1800 г.” - “Нужно тысяча восьмисотый год”

Rules could interfere:

- “Корпус.” - “корпустроение”

# Neural Normalizer



<https://storage.googleapis.com/pub-tools-public-publication-data/pdf/17a20b71a2c09100daaef5a8c39eeb930a7017f6.pdf>



# Neural Normalizer

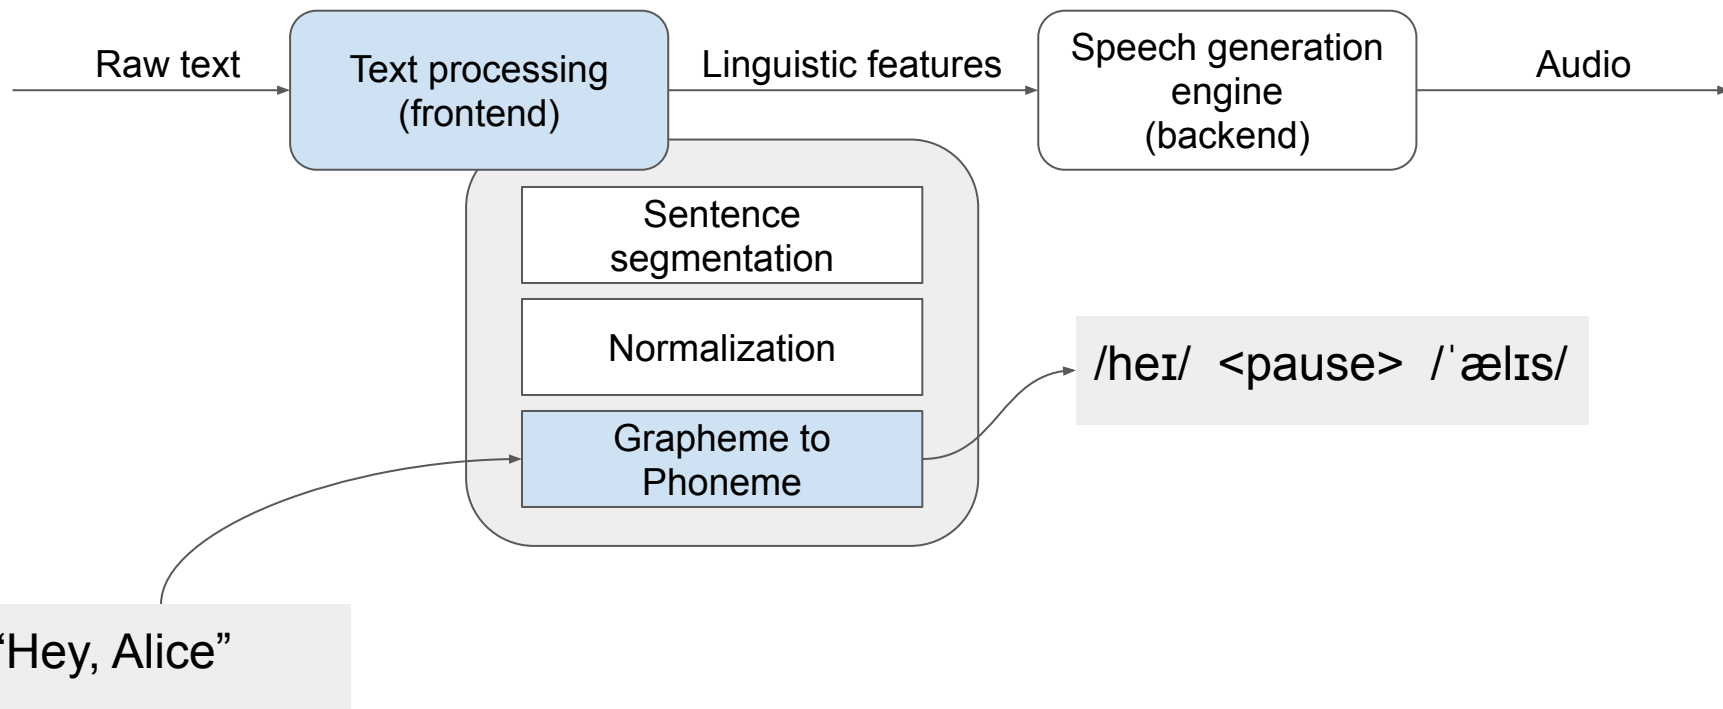
## Pros:

- Has no context problems
- No need to write code for new case, just add more data

## Cons:

- Need a lot of data ( $10^6$ - $10^7$  parallel sentences)
- Slower than WFST
- Impossible to fix intricate cases manually

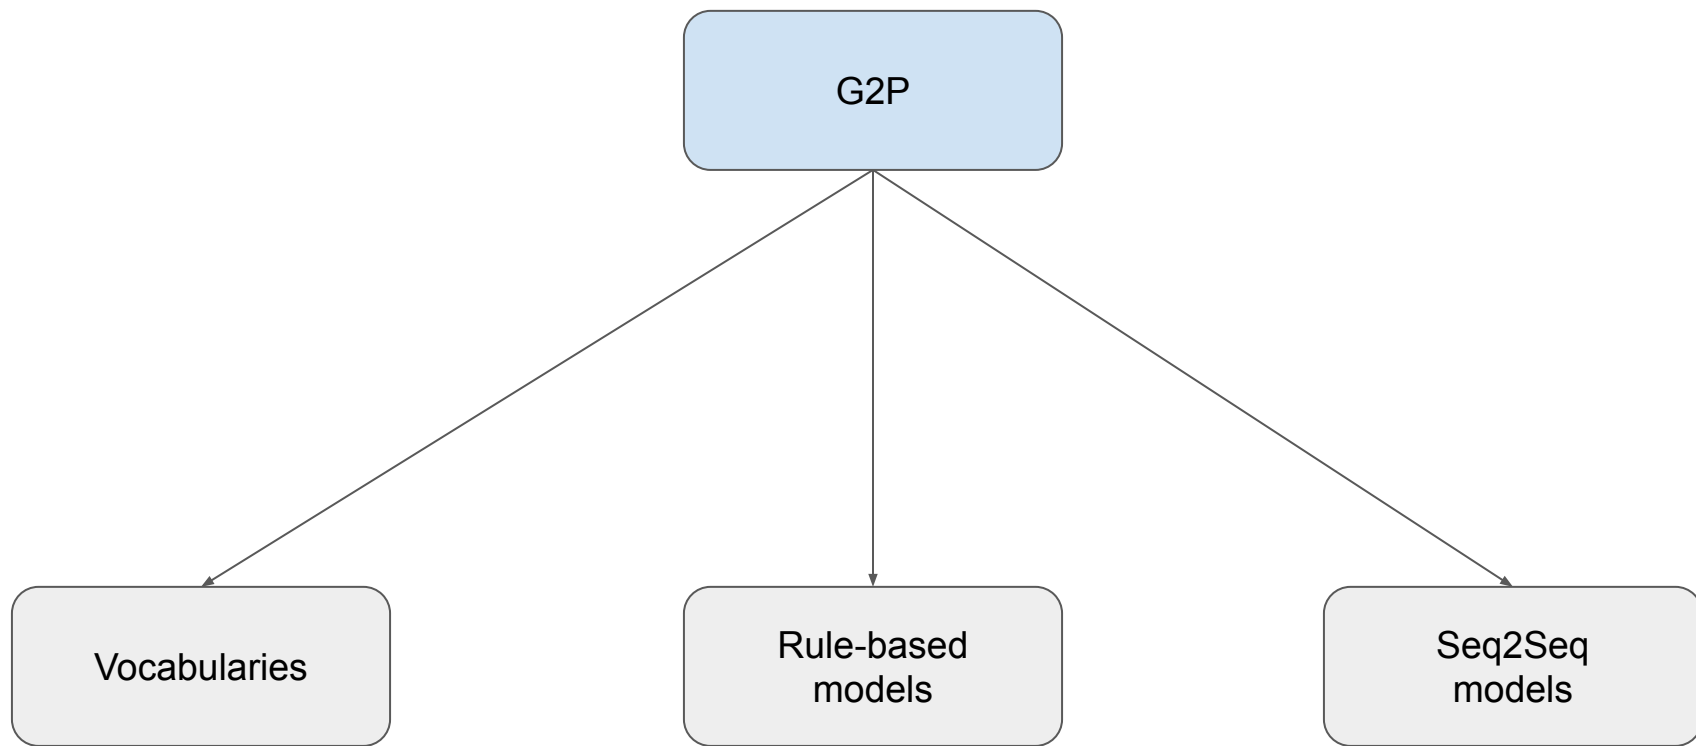
# Grapheme to Phoneme



# G2P Issues

- Homographs  
“Increase” - /ɪnˈkriːs/ or /ˈɪnkriːs/ ?
- Foreign words  
“A short poem **à la** Ogden Nash” - /ɑ la/
- Abbreviations  
“CI” - /siː/ /aɪ/

# G2P Components



# Phonemes or Texts?

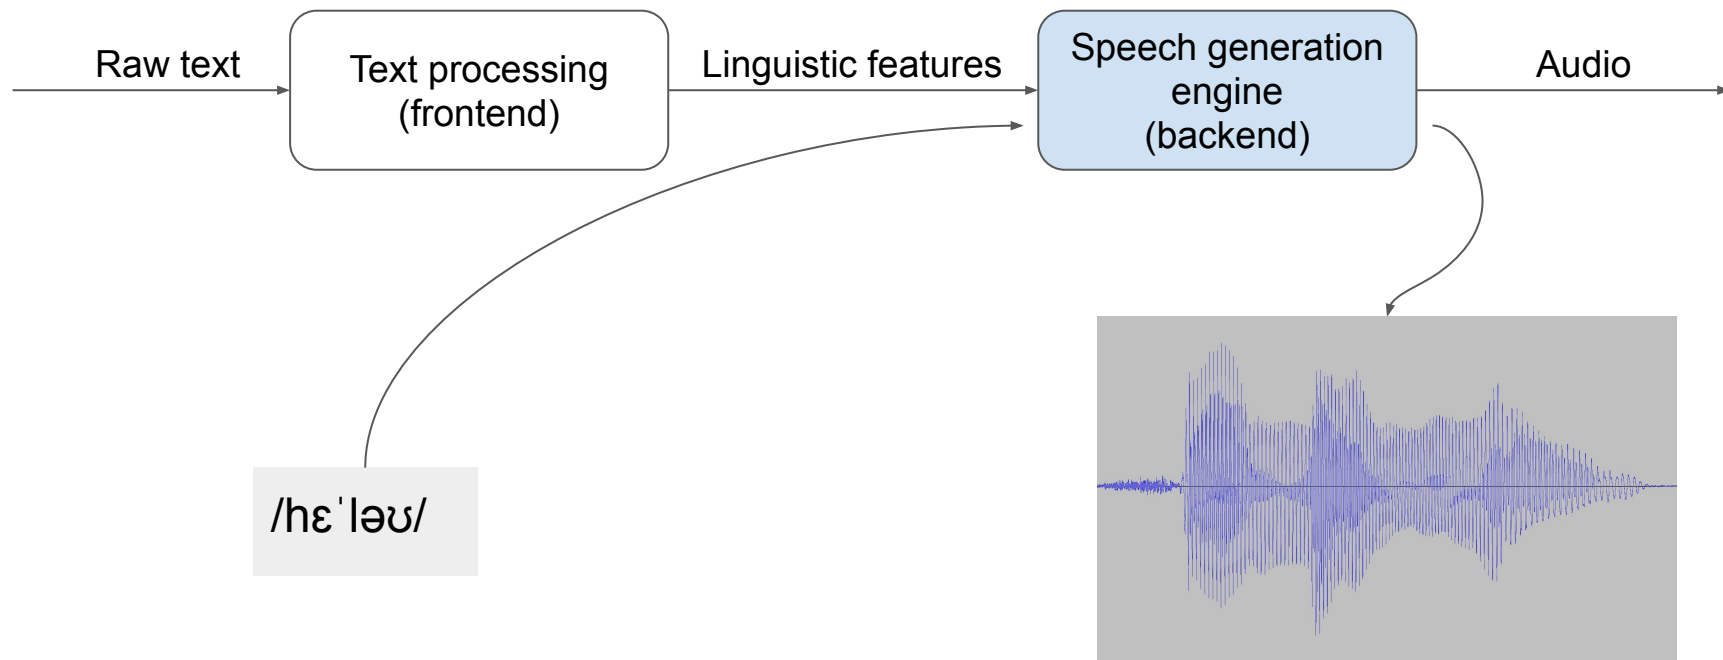
## G2P Pros:

- A separate G2P engine can serve for several backends
- Easy to control and fix certain words pronunciation

## G2P Cons:

- More expensive markup than text
- Additional components - additional bugs
- Additional information in text (e.g. punctuation) is lost
- Not all languages require G2P



# Engine



# Dense PCM

- Duration of phoneme (in Russian) ~ 20-150 ms
- Duration of 22.05kHz PCM sample ~ 0.045 ms

# Concatenative synthesis






# Unit selection

$\text{Join}(u_1, u_2)$

$\text{Target}(ph_4, u_4)$



$$\sum_i (\text{Join}(u_i, u_{i+1}) + \text{Target}(ph_i, u_i)) \rightarrow \min_{\{u\}}$$

\*Deep Learning for Siri Voice

<http://machinelearning.apple.com/2017/08/06/siri-voices.html>

# Unit selection

## Pros:

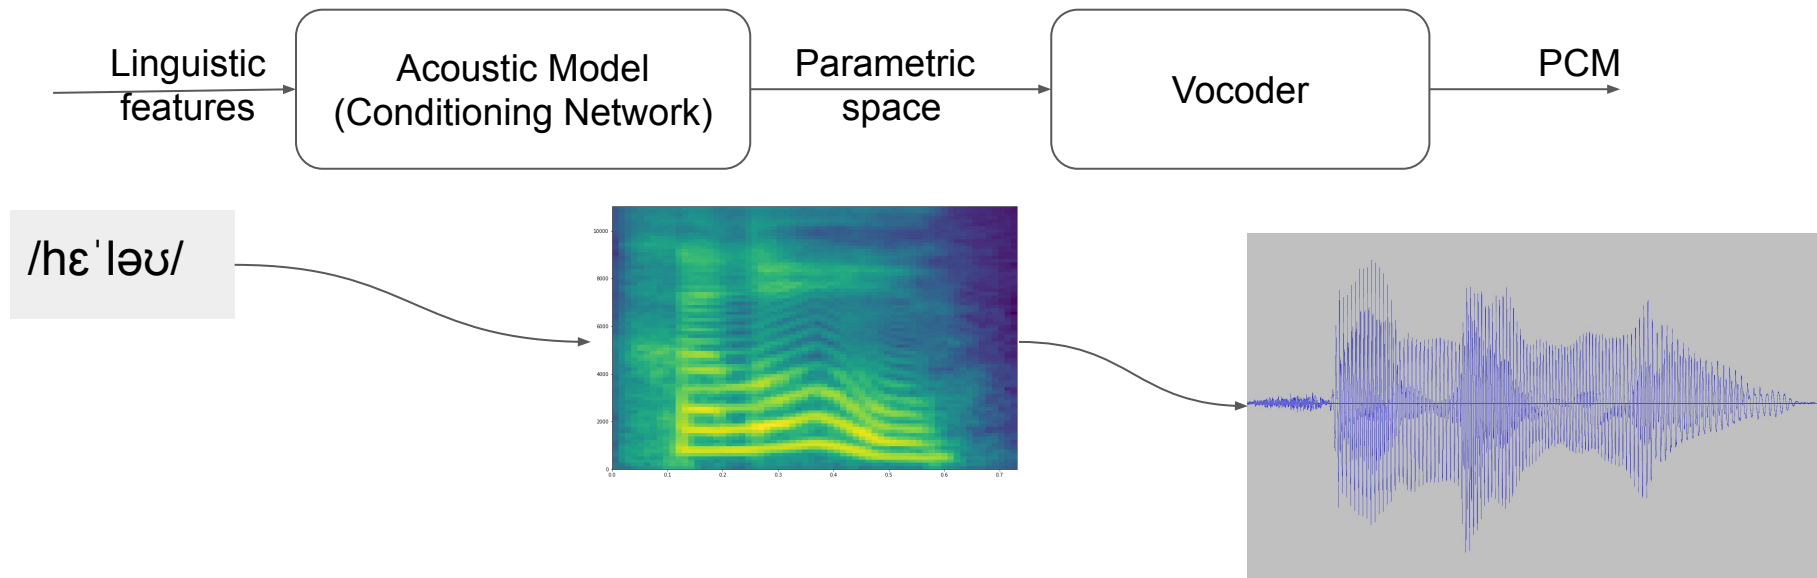
- Very fast synthesis, even with a high sample rate
- Sound quality inside units is equal to the original records

## Cons:

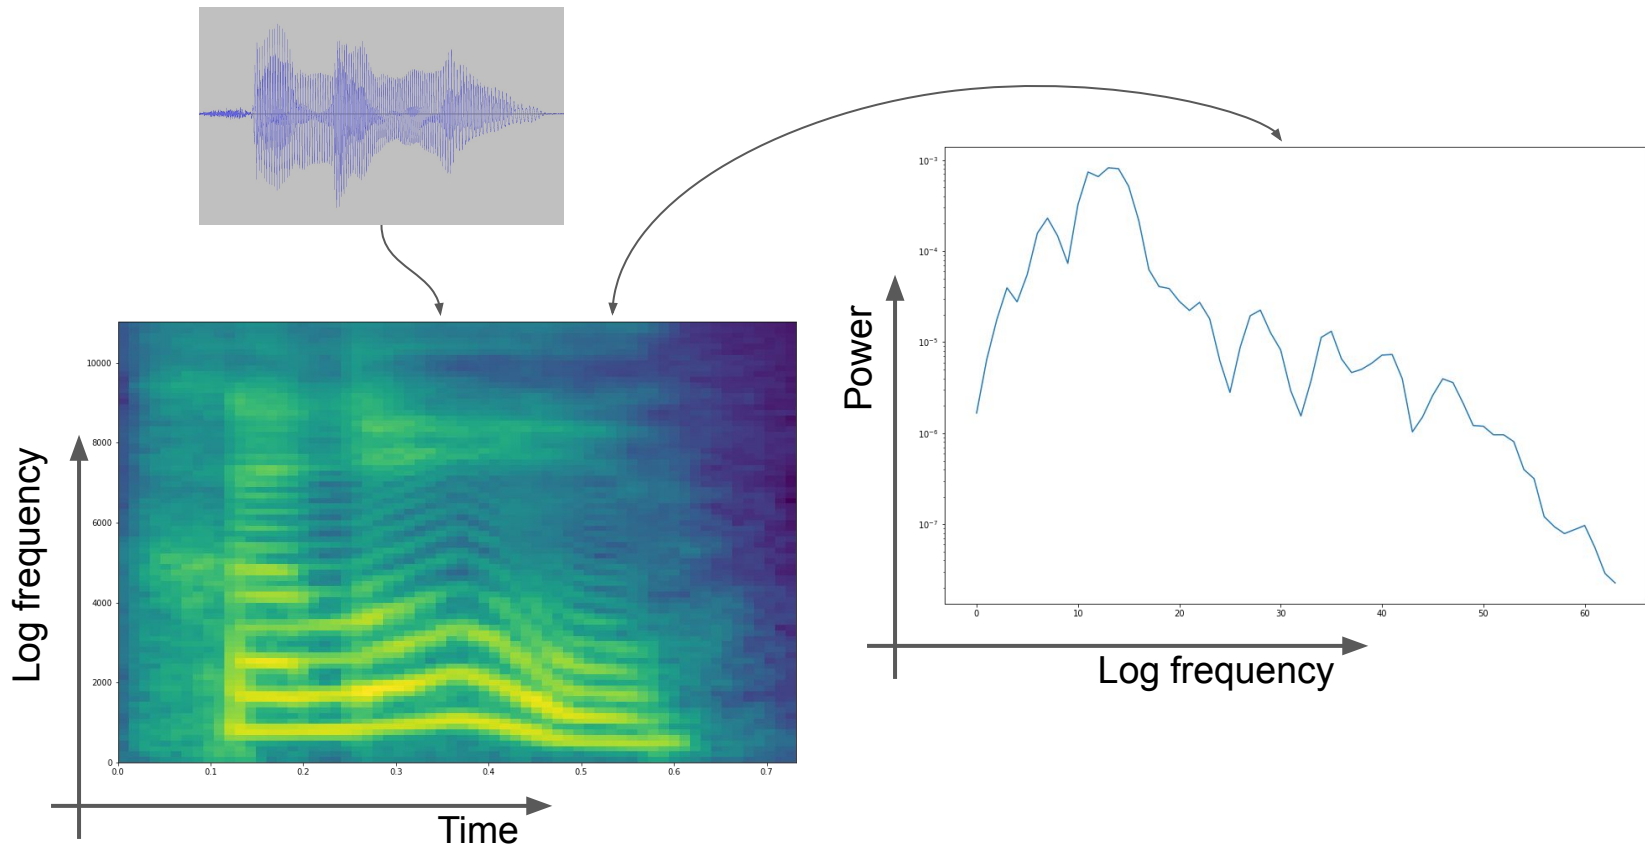
- Requires a lot of records to cover wide domain
- Produce annoying artifacts at bad unit concatenations
- Very bad at intonation variation
- Outside of records domain quality significantly decreases

# Parametric synthesis

- Introduce an intermediate speech audio representation
- Generate audio in two steps, by two models




# Parametric space - Mel Spectrogram



# Acoustic Model

- Deals with phoneme durations
- Solves regression problem for acoustic features
- Is responsible for intonation contour
- Is responsible for phoneme articulation
- (for models with text input) Solves G2P problem

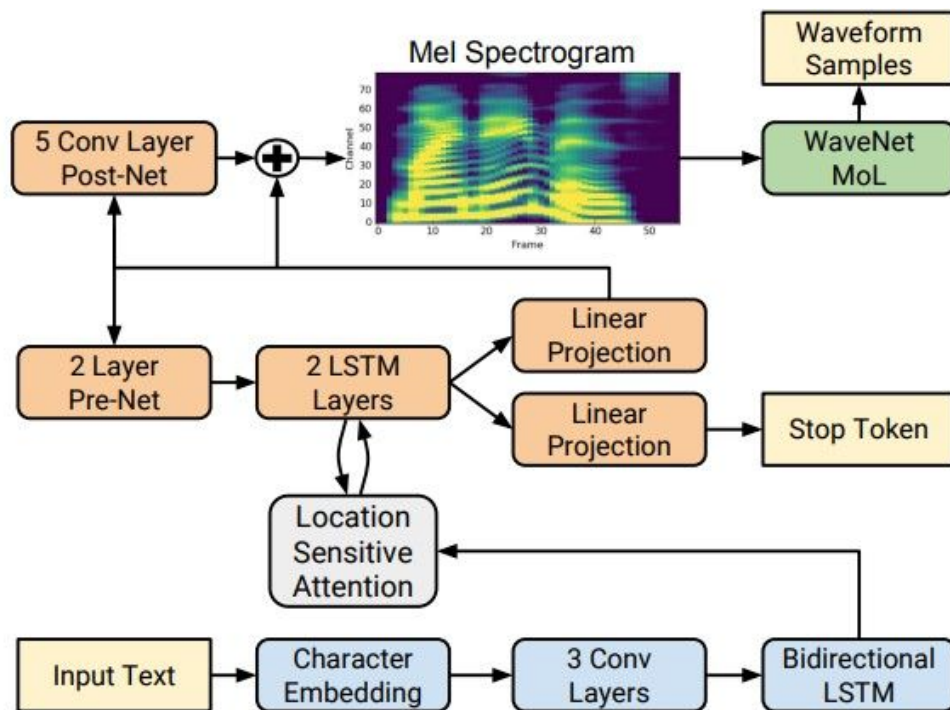


# Acoustic Model

Models used for acoustic modeling:

1. Decision Trees + HMM
2. Deep Mixture Density Networks + HMM
3. RNN with duration prediction  
Deep Voice 2 <https://arxiv.org/abs/1705.08947>
4. Seq2Seq Networks  
Char2Wav <https://openreview.net/pdf?id=B1VWyySKx>  
Deep Voice 3 <https://arxiv.org/abs/1710.07654>  
Tacotron 1 <https://arxiv.org/abs/1703.10135>

# Tacotron 2



**Fig. 1.** Block diagram of the Tacotron 2 system architecture.

\* Natural TTS Synthesis by Conditioning WaveNet on Mel Spectrogram Predictions  
<http://arxiv.org/pdf/1712.05884.pdf>

# Tacotron 2 - Train Loss

$$L_{pre} = ||melspec_{pre} - melspec_{gt}||_2^2$$

$$L_{post} = ||melspec_{post} - melspec_{gt}||_2^2$$

$$L_{stop} = \text{XEnt}(\text{stop token}, \mathbb{I}_{EOS})$$

$$L = L_{pre} + L_{post} + L_{stop}$$

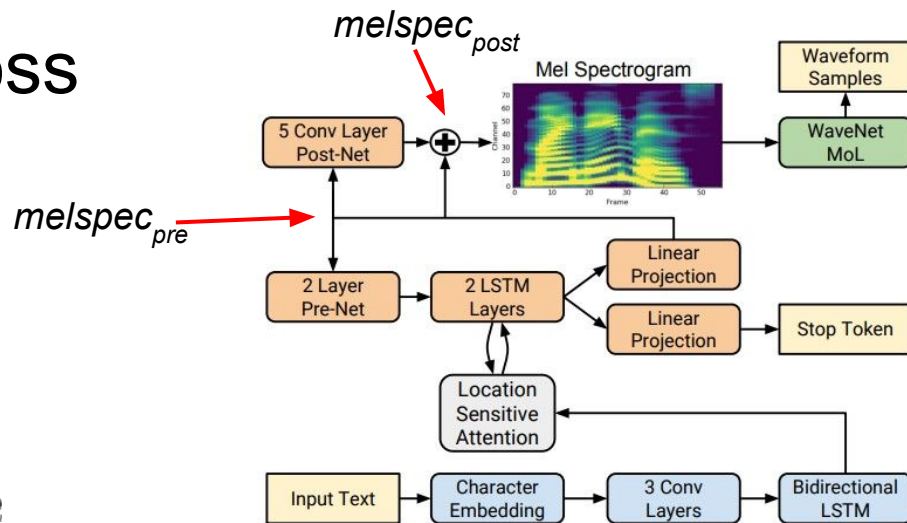


Fig. 1. Block diagram of the Tacotron 2 system architecture.

- No discrete output - cannot add EOS token
- Separate head for EOS prediction
- Two MSE losses for MelSpec prediction



# Tacotron 2 - Inference

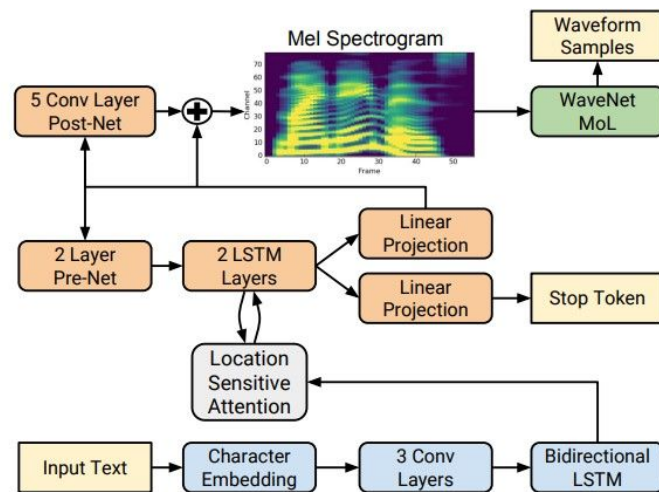
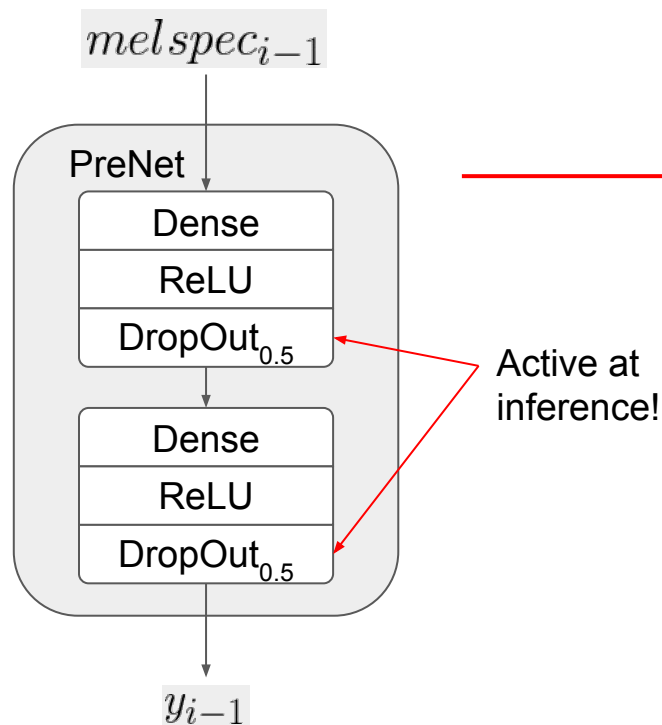
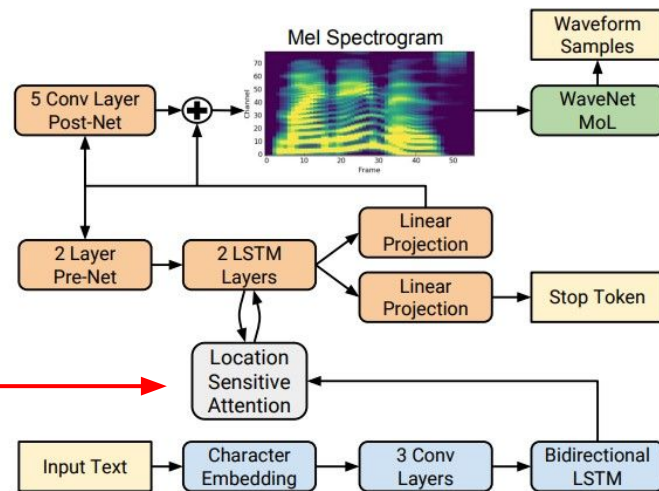



Fig. 1. Block diagram of the Tacotron 2 system architecture.

- $melspec_i$  is close to  $melspec_{i-1}$
- Need to prevent data leak in autoregression
- Otherwise the network learns to copy previous value

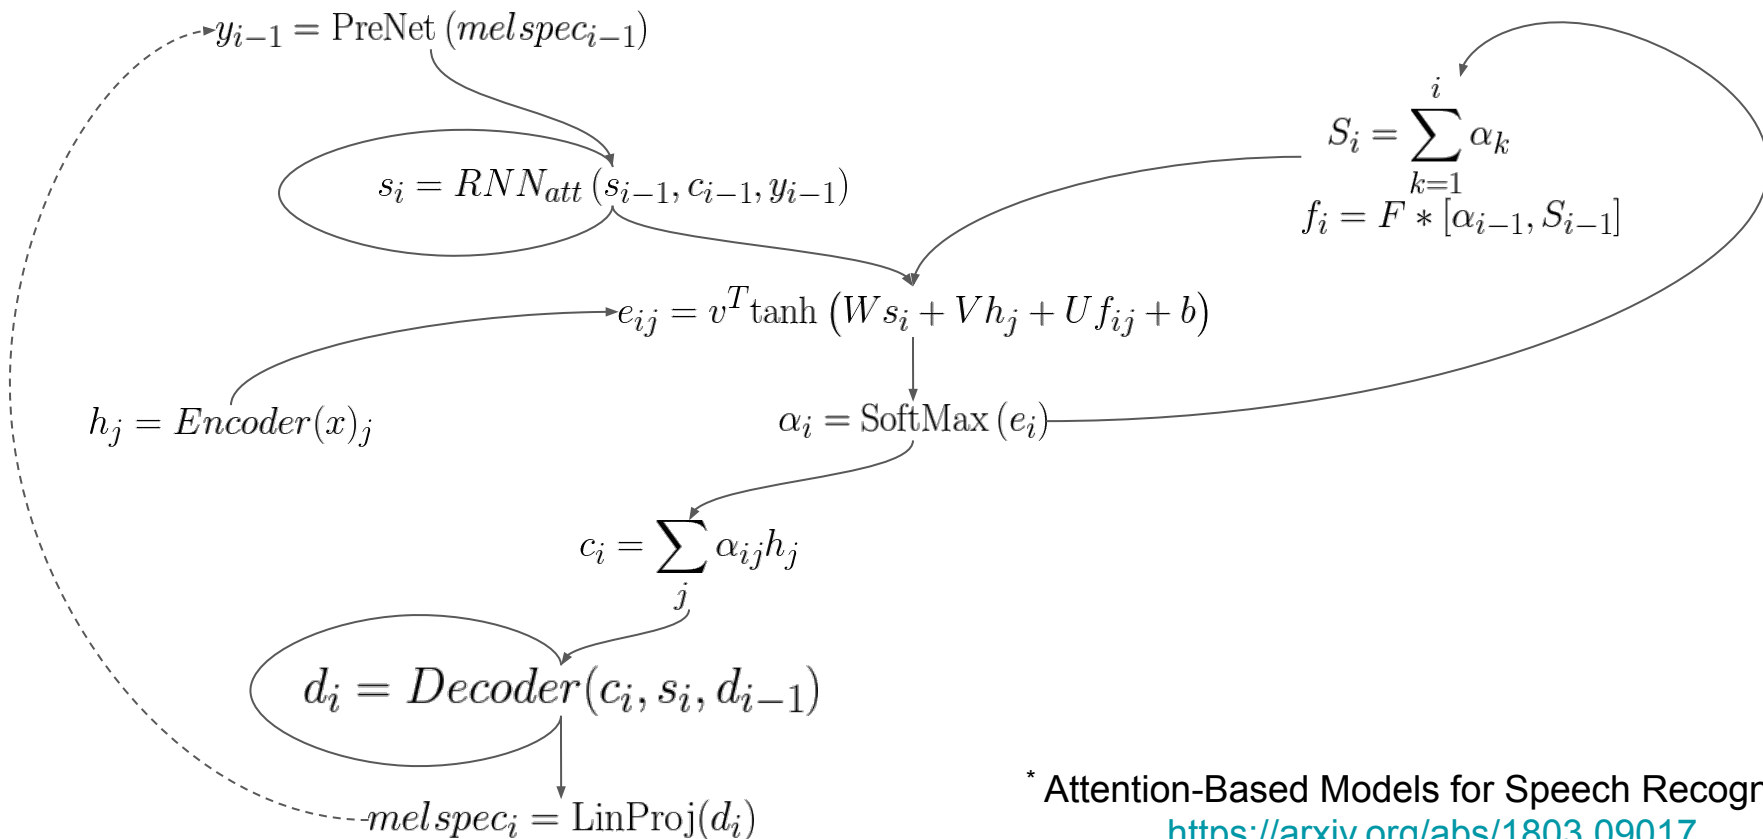
# Tacotron 2 - Attention

- Attention is (mostly) monotonic
- Depends on the previous state
- More stable than Bahdanau or Transformer attention

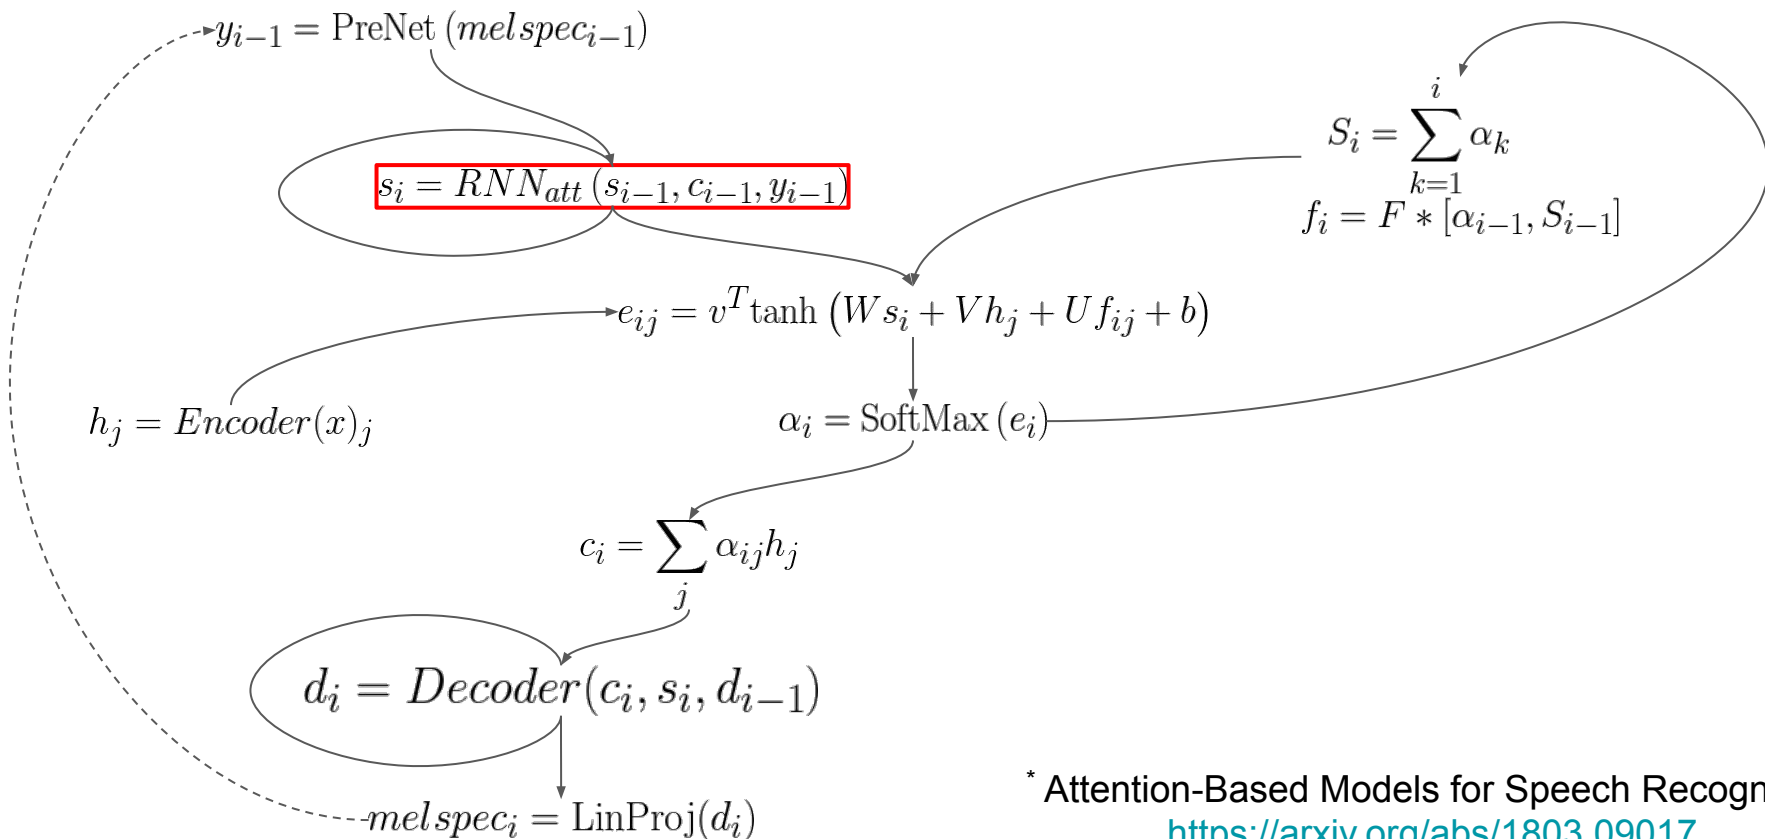


**Fig. 1.** Block diagram of the Tacotron 2 system architecture.

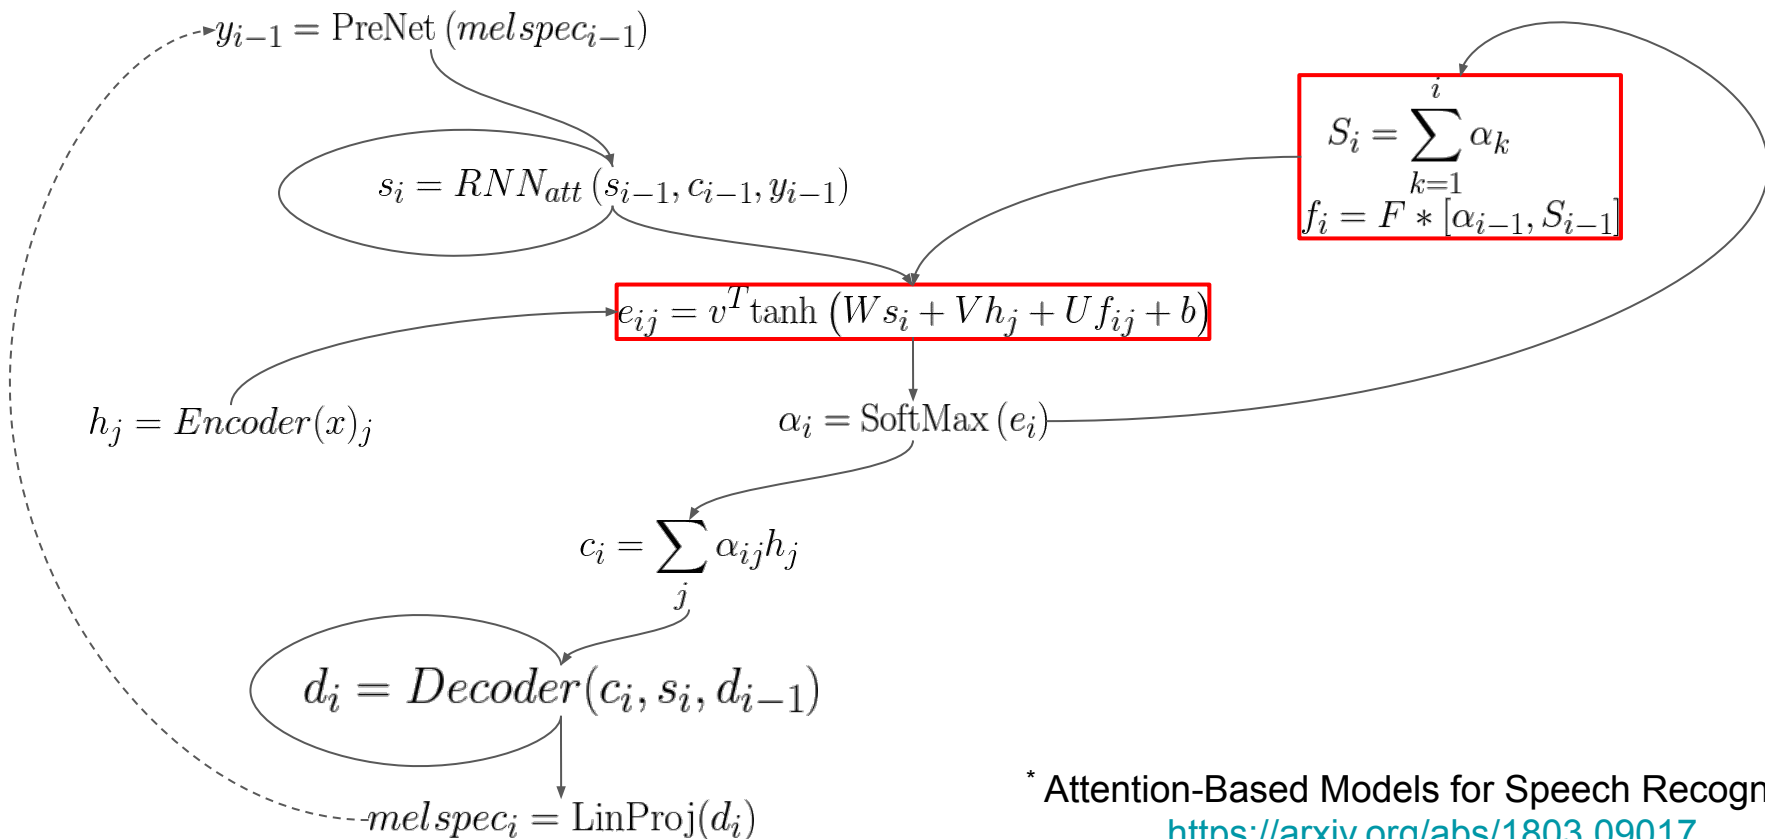
# Location-sensitive Attention



# Location-sensitive Attention



# Location-sensitive Attention



# What's Next? - Attention

More robust:

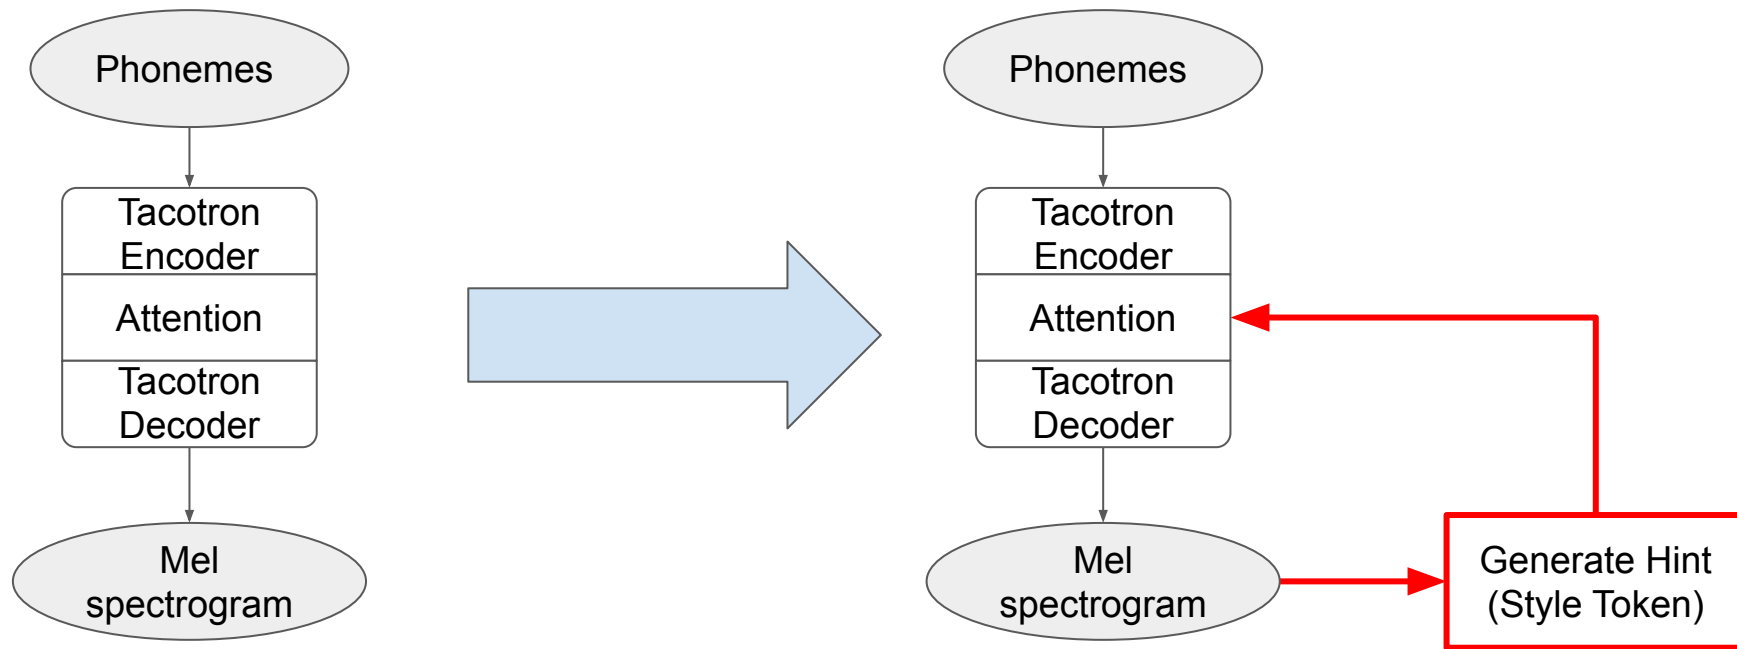
- Guided Attention <https://arxiv.org/abs/1710.08969>  
Stitch attention to record alignment (e.g. with ASR)
- Monotonic Attention <https://arxiv.org/abs/1906.00672>  
Enforce attention to read all input tokens sequentially
- Location-Relative Attention <https://arxiv.org/abs/1910.10288>  
Dynamic convolutions in attention make it more robust at very long utterances

# What's Next? - Attention

Faster:

- Transformer <https://arxiv.org/abs/1809.08895>  
Fast convergence  
Issues with monotonicity
- FastSpeech <https://arxiv.org/abs/1905.09263>  
Transformers with duration prediction (use pre-trained Transformer TTS)  
No autoregression - very fast inference  
Issues with quality

# What's Next? - Style Tokens



\* Style Tokens: Unsupervised Style Modeling, Control and Transfer in End-to-End Speech Synthesis  
<https://arxiv.org/abs/1803.09017>




# What's Next? - Style Tokens

- Different mechanisms of style modeling:
  - Several tokens with attention  
<https://arxiv.org/abs/1803.09017>
  - Variational Autoencoders  
<https://arxiv.org/abs/1804.02135>
- Predicting style:
  - From text encoder  
<https://arxiv.org/abs/1808.01410>
  - Use additional interpretable features  
<https://arxiv.org/abs/1810.07217>



# What's Next? - Multi-Speaker

- Multiple speaker training
  - Single model trained with multiple speakers/languages
  - More stable training
  - Transfer language/accents
  - Multilingual Multispeaker Tacotron <https://arxiv.org/abs/1907.04448>
- Speaker few-shot learning
  - Create new speaker for 1-2 records
  - Fast but unstable
  - Transfer Learning from Speaker Verification to Multispeaker <https://arxiv.org/abs/1806.04558>



# Vocoder

- Reconstructs audio from para-space
- Sets loudness, PCM sampling rate and precision
- Responsible for low-level sound quality



# Vocoder


- DSP Vocoder
  - Faster, smaller
  - Worse quality
  - Usually works with more complex para-space

World, Straight, Griffin-Lim

- Neural Vocoder
  - Requires much more computing power
  - Better quality (SotA)
  - Usually works with mel spectrograms

WaveNet, LPCNet, ClariNet, WaveRNN, WaveGlow, MelGAN


# WaveNet



\* WaveNet: A Generative Model for Raw Audio

<https://arxiv.org/abs/1609.03499>

# WaveNet - Output



- 8bit audio => classification in  $2^8$  classes - cannot work with 16bit
- To upsample  $\mu$ -law transformation was used

$$f(x_t) = \text{sign}(x_t) \frac{\ln(1 + \mu|x_t|)}{\ln(1 + \mu)}$$

- Later\* SML was replaced with a mixture of logistians:

$$\sum_i \pi_i \frac{e^{-(x-\mu_i)/s_i}}{s_i (1 + e^{-(x-\mu_i)/s_i})^2}$$

\* Improving the PixelCNN with Discretized Logistic Mixture Likelihood and Other Modifications  
<https://arxiv.org/abs/1701.05517>

# WaveNet

## Pros:


- Best audio quality for para-synthesis
- Works with continuous melspecs  
(good for MSE-loss Acoustic Model)

## Cons:

- **200** times slower than realtime

# Faster? Cache!

- No need to compute intermediate values twice
- WaveNet with caching is able to run in realtime



\* Fast Wavenet Generation Algorithm  
<https://arxiv.org/abs/1611.09482>



# WaveNet with Caching

## Pros:




- Realtime at 2-4 CPU cores
- Vanilla WaveNet quality

## Cons:


- Very hard to implement efficiently (assembler, special CPU instructions)
- Almost impossible to implement efficiently at GPU

# Faster? RNN!

- RNNs can keep the state
- No need to re-run the vocoder at each sample
- Smaller and faster networks



# WaveRNN



$$\mathbf{x}_t = [\mathbf{c}_{t-1}, \mathbf{f}_{t-1}, \mathbf{c}_t]$$

$$\mathbf{u}_t = \sigma(\mathbf{R}_u \mathbf{h}_{t-1} + \mathbf{I}_u^* \mathbf{x}_t)$$

$$\mathbf{r}_t = \sigma(\mathbf{R}_r \mathbf{h}_{t-1} + \mathbf{I}_r^* \mathbf{x}_t)$$

$$\mathbf{e}_t = \tau(\mathbf{r}_t \circ (\mathbf{R}_e \mathbf{h}_{t-1}) + \mathbf{I}_e^* \mathbf{x}_t)$$

$$\mathbf{h}_t = \mathbf{u}_t \circ \mathbf{h}_{t-1} + (1 - \mathbf{u}_t) \circ \mathbf{e}_t$$

$$\mathbf{y}_c, \mathbf{y}_f = \text{split}(\mathbf{h}_t)$$


$$P(\mathbf{c}_t) = \text{softmax}(\mathbf{O}_2 \text{relu}(\mathbf{O}_1 \mathbf{y}_c))$$

$$P(\mathbf{f}_t) = \text{softmax}(\mathbf{O}_4 \text{relu}(\mathbf{O}_3 \mathbf{y}_f))$$

\* Efficient Neural Audio Synthesis

<https://arxiv.org/abs/1802.08435v1>

# WaveRNN




$$\text{GRU}(\mathbf{x}, \mathbf{h}) \sim \begin{cases} \mathbf{x}_t = [\mathbf{c}_{t-1}, \mathbf{f}_{t-1}, \mathbf{c}_t] \\ \mathbf{u}_t = \sigma(\mathbf{R}_u \mathbf{h}_{t-1} + \mathbf{I}_u^* \mathbf{x}_t) \\ \mathbf{r}_t = \sigma(\mathbf{R}_r \mathbf{h}_{t-1} + \mathbf{I}_r^* \mathbf{x}_t) \\ \mathbf{e}_t = \tau(\mathbf{r}_t \circ (\mathbf{R}_e \mathbf{h}_{t-1}) + \mathbf{I}_e^* \mathbf{x}_t) \\ \mathbf{h}_t = \mathbf{u}_t \circ \mathbf{h}_{t-1} + (1 - \mathbf{u}_t) \circ \mathbf{e}_t \\ \mathbf{y}_c, \mathbf{y}_f = \text{split}(\mathbf{h}_t) \\ P(\mathbf{c}_t) = \text{softmax}(\mathbf{O}_2 \text{relu}(\mathbf{O}_1 \mathbf{y}_c)) \\ P(\mathbf{f}_t) = \text{softmax}(\mathbf{O}_4 \text{relu}(\mathbf{O}_3 \mathbf{y}_f)) \end{cases}$$

\* Efficient Neural Audio Synthesis

<https://arxiv.org/abs/1802.08435v1>


# WaveRNN



$$\text{GRU}(x, h) \sim \begin{cases} \mathbf{x}_t = [\mathbf{c}_{t-1}, \mathbf{f}_{t-1}, \mathbf{c}_t] \\ \mathbf{u}_t = \sigma(\mathbf{R}_u \mathbf{h}_{t-1} + \mathbf{I}_u^* \mathbf{x}_t) \\ \mathbf{r}_t = \sigma(\mathbf{R}_r \mathbf{h}_{t-1} + \mathbf{I}_r^* \mathbf{x}_t) \\ \mathbf{e}_t = \tau(\mathbf{r}_t \circ (\mathbf{R}_e \mathbf{h}_{t-1}) + \mathbf{I}_e^* \mathbf{x}_t) \\ \mathbf{h}_t = \mathbf{u}_t \circ \mathbf{h}_{t-1} + (1 - \mathbf{u}_t) \circ \mathbf{e}_t \\ \mathbf{y}_c, \mathbf{y}_f = \text{split}(\mathbf{h}_t) \\ P(\mathbf{c}_t) = \text{softmax}(\mathbf{O}_2 \text{relu}(\mathbf{O}_1 \mathbf{y}_c)) \\ P(\mathbf{f}_t) = \text{softmax}(\mathbf{O}_4 \text{relu}(\mathbf{O}_3 \mathbf{y}_f)) \end{cases}$$

$$1. (h_{t-1}, [c_{t-1}, c_{t-1}]) \rightarrow h_t^c \rightarrow y_c \rightarrow c_t$$

# WaveRNN



$$\text{GRU}(x, h) \sim \begin{cases} x_t = [c_{t-1}, f_{t-1}, c_t] \\ u_t = \sigma(R_u h_{t-1} + I_u^* x_t) \\ r_t = \sigma(R_r h_{t-1} + I_r^* x_t) \\ e_t = \tau(r_t \circ (R_e h_{t-1}) + I_e^* x_t) \\ h_t = u_t \circ h_{t-1} + (1 - u_t) \circ e_t \\ y_c, y_f = \text{split}(h_t) \\ P(c_t) = \text{softmax}(O_2 \text{relu}(O_1 y_c)) \\ P(f_t) = \text{softmax}(O_4 \text{relu}(O_3 y_f)) \end{cases}$$

1.  $(h_{t-1}, [c_{t-1}, c_{t-1}]) \rightarrow h_t^c \rightarrow y_c \rightarrow c_t$
2.  $(h_{t-1}, [c_{t-1}, f_{t-1}, c_t]) \rightarrow h_t^f \rightarrow y_f \rightarrow f_t$

# WaveRNN

## Pros:




- Real-time inference
- Quality comparable with WaveNet
- Can run (compressed version) at CPU in realtime

## Cons:

- Hard to make GPU inference effective

# Faster? Parallel!

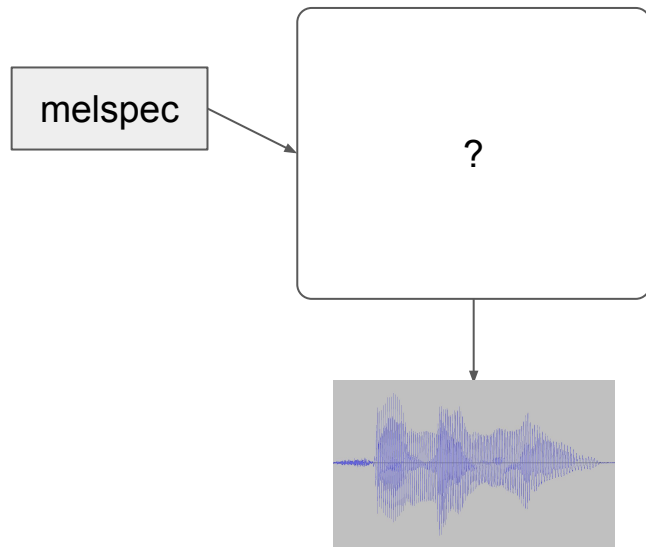
- Very efficient generation at GPU
- Need to break autoregression loop
- Vocoder:
  - Parallel WaveNet <https://arxiv.org/abs/1711.10433>
  - ClariNet <https://arxiv.org/abs/1807.07281>
  - WaveGlow <https://arxiv.org/abs/1811.00002>





# WaveGlow

1. Need to sample:  $x \sim D_{wav}(\text{melspec})$
2.  $D_{wav} - ?$

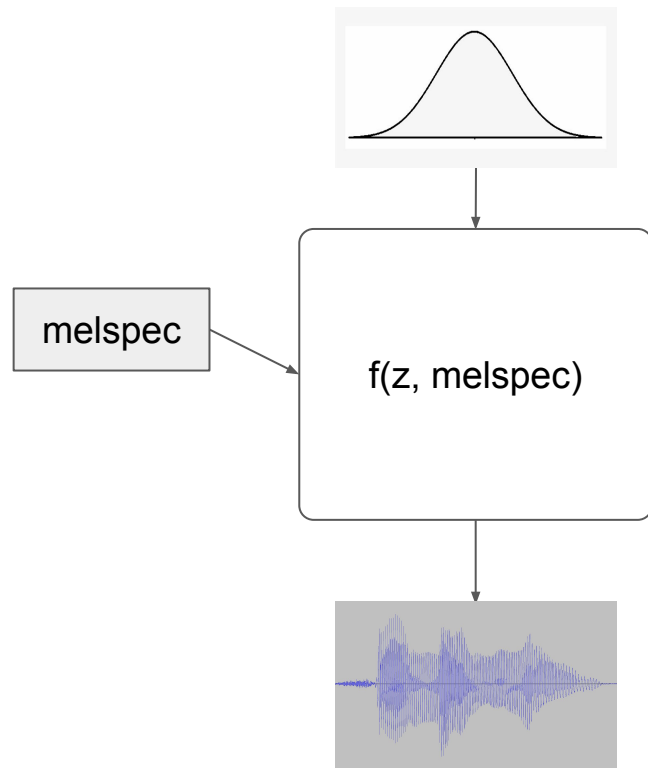


\* WaveGlow: A Flow-based Generative Network for Speech Synthesis

<https://arxiv.org/abs/1811.00002>

# WaveGlow

1. Need to sample:  $x \sim D_{wav}(\text{melspec})$
2.  $D_{wav} - ?$
3.  $x = f(z, \text{melspec}); z \sim \mathcal{N}(0, 1)$



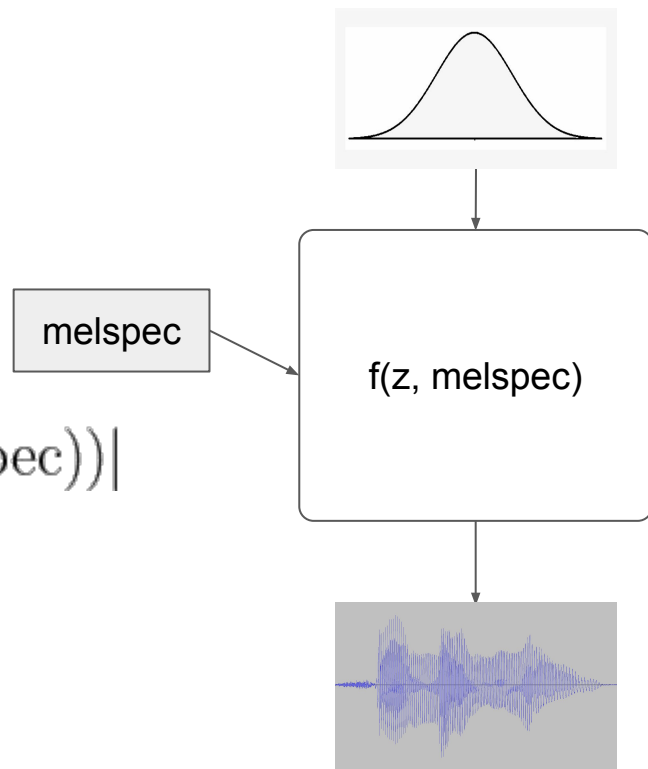
\* WaveGlow: A Flow-based Generative Network for Speech Synthesis

<https://arxiv.org/abs/1811.00002>

# WaveGlow

1. Need to sample:  $x \sim D_{wav}(\text{melspec})$
2.  $D_{wav} - ?$
3.  $x = f(z, \text{melspec}); z \sim \mathcal{N}(0, 1)$
4.  $z = f^{-1}(x)$

$$p(x|\text{melspec}) = p(z(x)) |\det J(f^{-1}(x, \text{melspec}))|$$



\* WaveGlow: A Flow-based Generative Network for Speech Synthesis

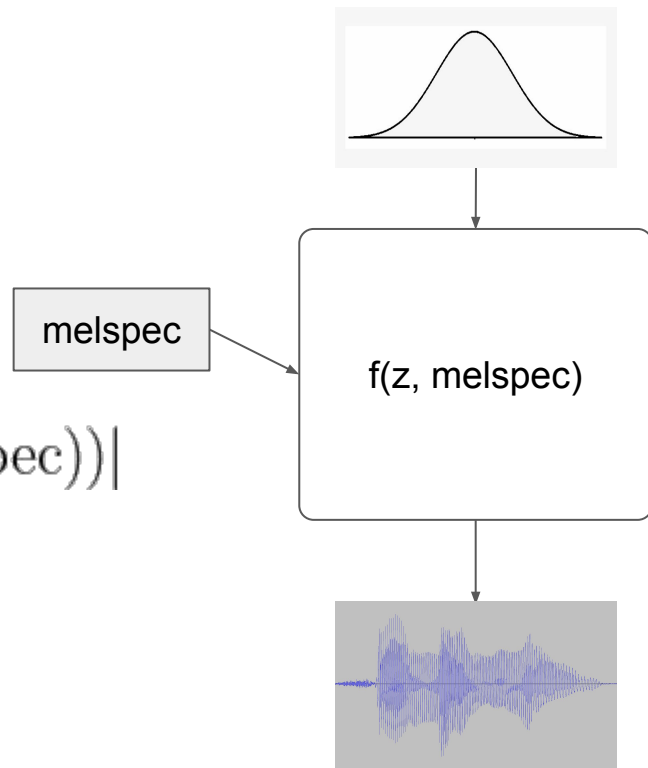
<https://arxiv.org/abs/1811.00002>

# WaveGlow

1. Need to sample:  $x \sim D_{wav}(\text{melspec})$
2.  $D_{wav} - ?$
3.  $x = f(z, \text{melspec}); z \sim \mathcal{N}(0, 1)$
4.  $z = f^{-1}(x)$

$$p(x|\text{melspec}) = p(z(x)) |\det J(f^{-1}(x, \text{melspec}))|$$

5. Train to maximize  $p_{\theta}(x|\text{melspec})$



\* WaveGlow: A Flow-based Generative Network for Speech Synthesis

<https://arxiv.org/abs/1811.00002>

# WaveGlow

1. Need to sample:  $x \sim D_{wav}(\text{melspec})$

2.  $D_{wav} - ?$

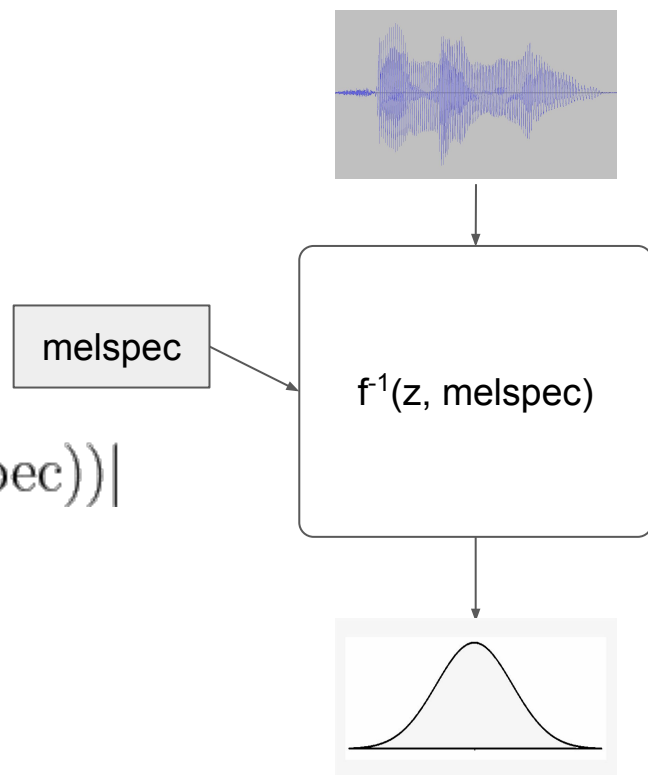
3.  $x = f(z, \text{melspec}); z \sim \mathcal{N}(0, 1)$

4.  $z = f^{-1}(x)$

$$p(x|\text{melspec}) = p(z(x)) |\det J(f^{-1}(x, \text{melspec}))|$$

5. Train to maximize  $p_{\theta}(x|\text{melspec})$

6. Train invertible function  $f^{-1}(x, \text{melspec})$



\* WaveGlow: A Flow-based Generative Network for Speech Synthesis

<https://arxiv.org/abs/1811.00002>

# WaveGlow

1. Need to sample:  $x \sim D_{wav}(\text{melspec})$

2.  $D_{wav} - ?$

3.  $x = f(z, \text{melspec}); z \sim \mathcal{N}(0, 1)$

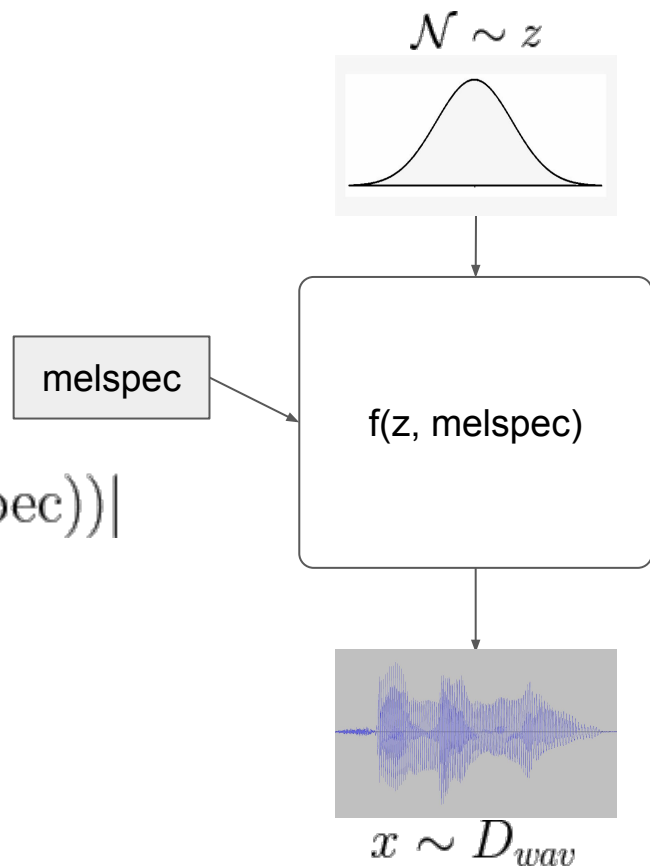
4.  $z = f^{-1}(x)$

$$p(x|\text{melspec}) = p(z(x)) |\det J(f^{-1}(x, \text{melspec}))|$$

5. Train to maximize  $p_{\theta}(x|\text{melspec})$

6. Train invertible function  $f^{-1}(x, \text{melspec})$

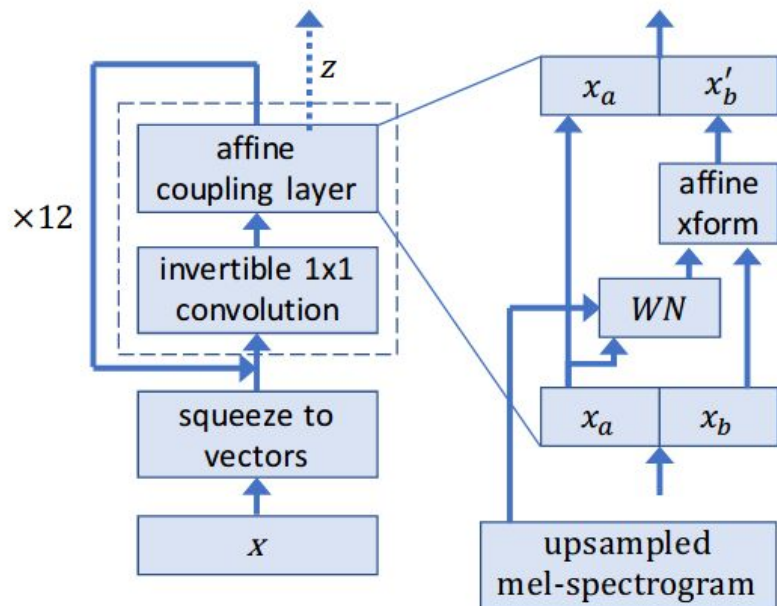
7. Revert and sample  $x$  via  $z$



\* WaveGlow: A Flow-based Generative Network for Speech Synthesis

<https://arxiv.org/abs/1811.00002>

# WaveGlow - Coupling



**Fig. 1:** WaveGlow network

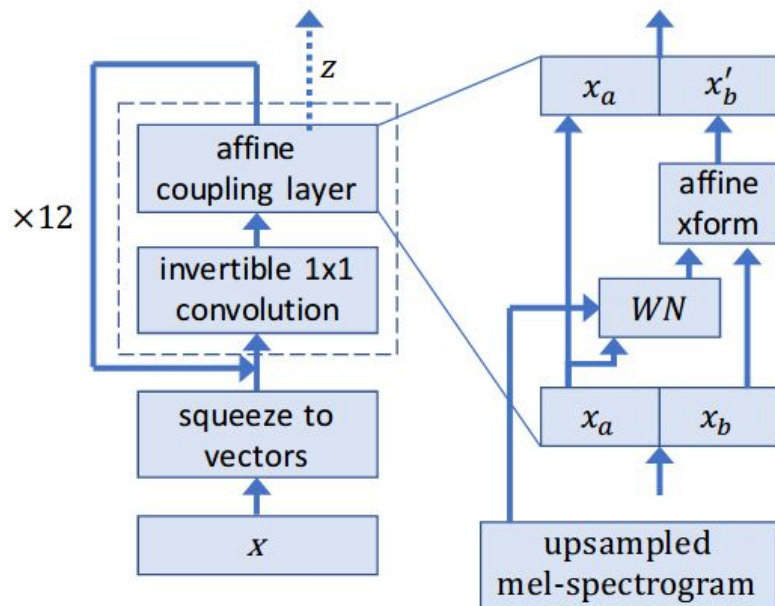
$$x_a, x_b = \text{split}(x)$$

$$(\log s, t) = WN(x_a, \text{melspec})$$

$$x'_b = s \odot x_b + t$$

$$f_{\text{coupling}}^{-1}(x) = \text{concat}(x_a, x'_b)$$

# WaveGlow - Coupling



**Fig. 1:** WaveGlow network

$$\begin{aligned}
 x_a, x_b &= \text{split}(x) \\
 (\log s, t) &= WN(x_a, \text{melspec}) \\
 x'_b &= s \odot x_b + t \\
 f_{\text{coupling}}^{-1}(x) &= \text{concat}(x_a, x'_b)
 \end{aligned}$$



$$\begin{aligned}
 x'_a, x'_b &= \text{split}(x') \\
 (\log s, t) &= WN(x'_a, \text{melspec}) \\
 x_b &= s^{-1} \odot (x'_b - t) \\
 f_{\text{coupling}}(x) &= \text{concat}(x'_a, x_b)
 \end{aligned}$$





# WaveGlow

## Pros:


- Quality comparable to vanilla WaveNet
- Fast (520k samples / second) inference at GPU
- Code at GitHub

## Cons:

- Difficult streaming inference
- Impossible to implement on CPU

# Hack of the Day

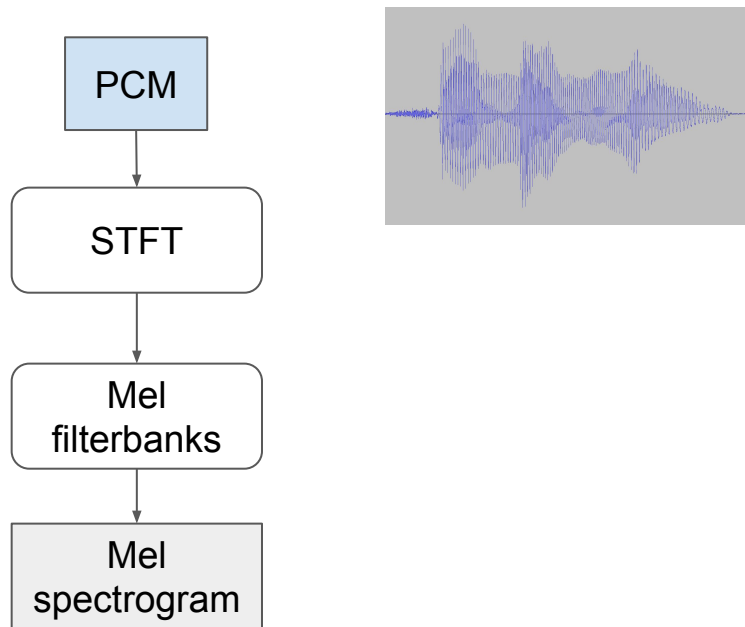
1. There is a github repo for it
  - <https://github.com/NVIDIA/waveglow>
  - <https://github.com/NVIDIA/tacotron2>
  - [https://github.com/r9y9/wavenet\\_vocoder](https://github.com/r9y9/wavenet_vocoder)
2.  $P(\text{"there's a bug"}) \gg P(\text{"it underfits"})$
3. Don't trust losses. Listen



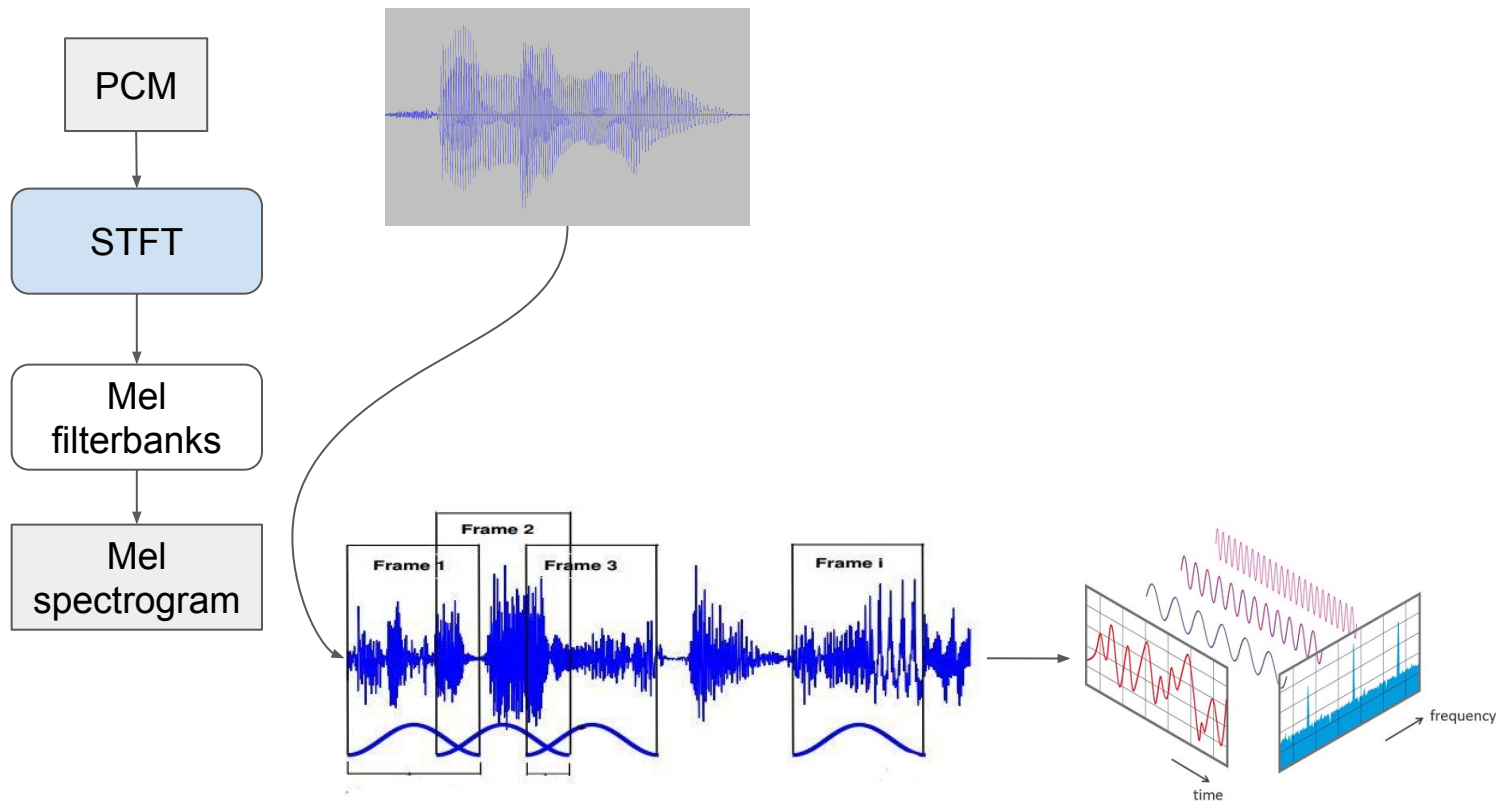
?

Extras

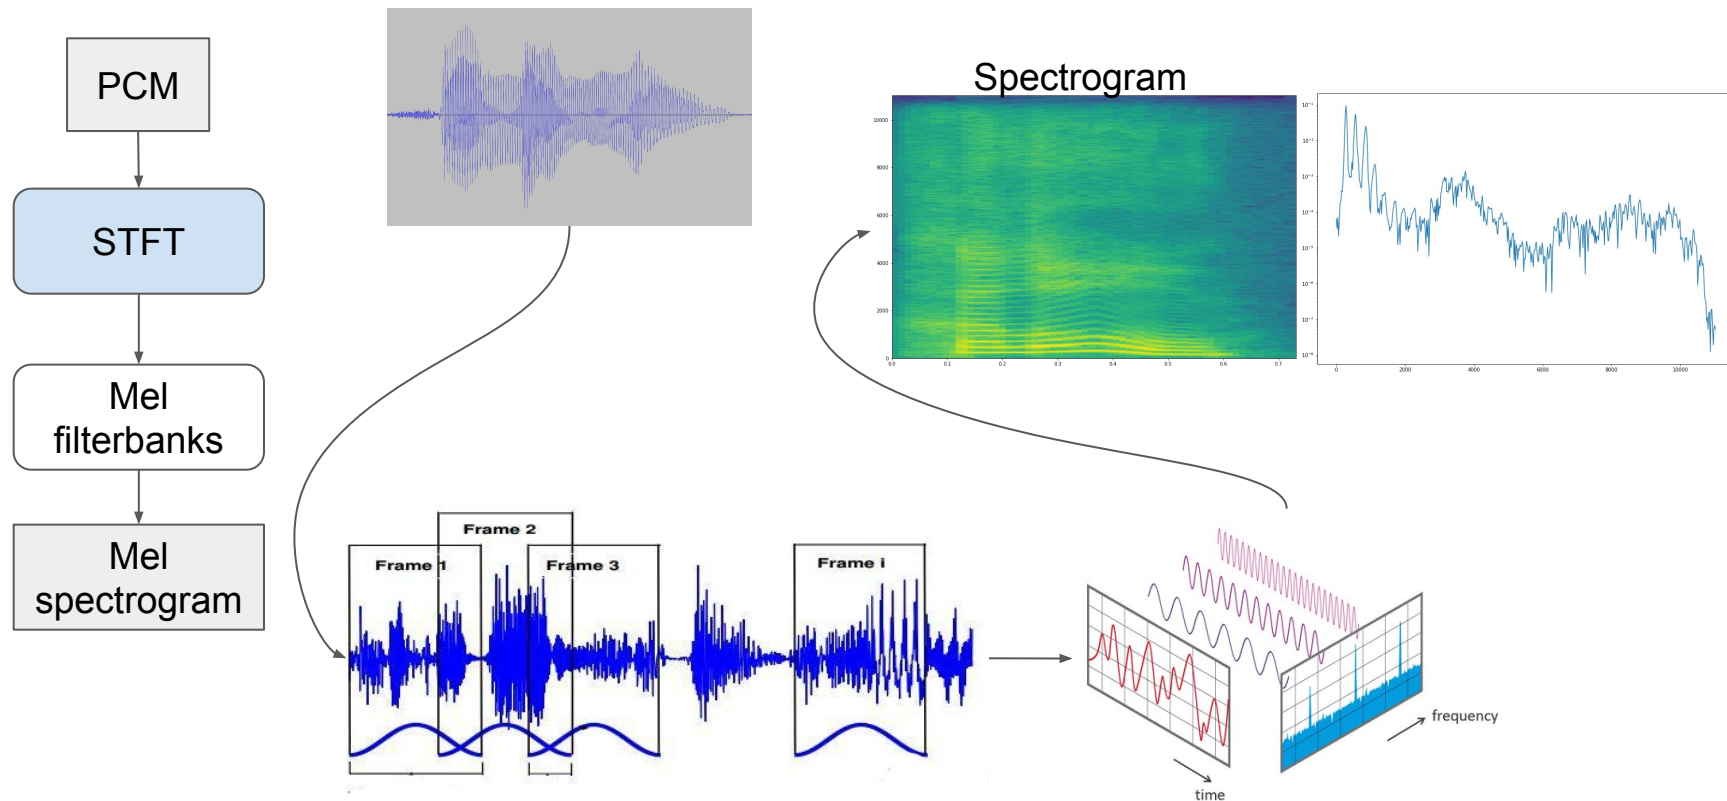
# Parametric space - Mel Spectrogram



# Parametric space - Mel Spectrogram





# Parametric space - Mel Spectrogram









# Parametric space - Mel Spectrogram





# Parametric space - Mel Spectrogram





Spectrogram




# Parametric space - Mel Spectrogram




Spectrogram



Mel Spectrogram



# Other para-spaces



# Other para-spaces

We could further process MelSpec to make more complex and compact features:

- F0 - frequency at which vocal chords vibrate in voiced sound
- Spectral Envelope, could be described with:
  - Cepstral Coefficients - Discrete Cosine Transform of (mel) spectrum
  - Linear Predictive Coding - autoregressive model describing envelope
- Periodicity
  - Frequency-wise
  - Max Voiced Frequency